

Towards Exascale Multiparticle Simulations

Srinath Kailasa

A thesis submitted in partial fulfillment of the requirements
for the degree Master of Philosophy

Department of Mathematics
University College London
November, 2022

Declaration

I, Srinath Kailasa, confirm that the work presented in this thesis is my own. Where information has been derived from other sources, I confirm that this has been indicated in the thesis.

Abstract

The past three decades have seen the emergence of so called ‘fast algorithms’ that are able to optimally apply and invert dense matrices that exhibit a special low-rank structure in their off-diagonal elements. Such matrices arise in numerous areas of science and engineering, for example in the linear system matrices of boundary integral formulations of problems from acoustics and electromagnetics to fluid dynamics, geomechanics and even seismology. In the best case matrices can be stored, applied and inverted in $O(N)$, in contrast to $O(N^2)$ for storage and application, and $O(N^3)$ for inversion when computed naively.

The unification of software for the forward and inverse application of these operators in a single set of open-source libraries optimised for distributed computing environments is lacking, and is the central concern of this research project. We propose the creation of a unified solver infrastructure that can demonstrate good weak scaling from local workstations to upcoming exascale machines. Developing high-performance implementations of fast algorithms is challenging due to highly-technical nature of their underlying mathematical machinery, further complicated by the diversity of software and hardware environments in which research code is expected to run.

This subsidiary thesis presents current progress towards this goal. Chapter (1) introduces the Fast Multipole Method (FMM), the prototypical fast algorithm for $O(N)$ matrix vector products, and discusses implementation strategies in the context of high-performance software implementations. Chapter (2) provides a survey of the fragmented parallel software landscape for fast algorithms, before proceeding with a case study of a Python implementation of an FMM, which attempted to bridge the gap between a familiar and ergonomic language for researchers and achieving high-performance. The remainder of the chapter introduces Rust, our proposed solution for ergonomic and high-performance codes for computational science, and it concludes with an overview of a software output: Rusty Tree, a new Rust-based library for the construction of parallel octrees, a foundational datastructure for FMMs, as well as other fast algorithms. Chapter (3) introduces vectors for future research, specifically an introduction to fast algorithms and software for matrix inversion, and the potential pitfalls we will face in their implementation for performance, as well as an overview of a proposed investigation into the optimal mathematical implementation of field translations - a crucial component of a performant FMM. We conclude with a look ahead towards a key target application for our software, the solution of electromagnetic scattering problems described by Maxwell’s equations that can demonstrate performance at exascale.

Contents

1	Introduction	1
1.1	Motivation	1
1.2	From Analytic to Algebraic Hierarchical Fast Multipole Methods . . .	3
2	Designing Software for Fast Algorithms	12
2.1	The Software Landscape	12
2.2	Case Study: PyExaFMM, a Python Fast Multipole Method	15
2.3	Rust for High Performance Scientific Computing	25
2.4	Case Study: RustyTree, a Rust based Parallel Octree	34
3	Looking Ahead	48
3.1	Fast Direct Solvers on Distributed Memory Systems	48
3.1.1	\mathcal{T}	49
3.1.2	\mathcal{K}'	51
3.1.3	\mathcal{S}	52
3.1.4	\mathcal{D}	53
3.2	Example Problems	53
3.2.1	Acoustic Sound Hard Scattering	53
3.2.2	Acoustic Transmission	55
3.3	Optimal Translation Operators for Fast Algorithms	56
3.4	Target Application: Maxwell Scattering	56
4	Conclusion	57
A	Appendix	58
A.1	Deriving Local Expansion Coefficients from Multipole Expansion in \mathbb{R}^2	58
A.2	Complexity of Tree Construction	58
A.3	HykSort	58
	Bibliography	59

Introduction

1.1 Motivation

The motivation behind the development of the original fast multipole method (FMM), was the calculation of potentials in N -body problems,

$$\phi_j = \sum_{i=1}^N K(x_i, x_j) q_i \quad (1.1)$$

Consider electrostatics, or gravitation, where q_i is a point charge or mass, and the kernels are of the form $K(x, y) = \log|x - y|$ in \mathbb{R}^2 , or $K(x, y) = \frac{1}{|x-y|}$ in \mathbb{R}^3 . Similar sums appear in the discretised form of boundary integral equation (BIE) formulations for elliptic partial differential equations (PDEs), which are the example that motivates our research. Generically, an integral equation formulation can be written as,

$$a(x)u(x) + b(x) \int_{\Omega} K(x, y)c(y)u(y)dy = f(x), \quad x \in \Omega \subset \mathbb{R}^d \quad (1.2)$$

where the dimension $d = 2$ or 3 . The functions $a(x)$, $b(x)$ and $c(y)$ are given and linked to the parameters of a problem, $K(x, y)$ is some known kernel function and $f(x)$ is a known right hand side, $K(x, y)$ is associated with the PDE - either its Green's function, or the derivative. This is a very general formulation, and includes common problems such as the Laplace and Helmholtz equations. Upon discretisation with an appropriate method, for example the Nyström or Galerkin methods, we obtain a linear system of the form,

$$\mathbf{K}u = f \quad (1.3)$$

The key feature of this linear system is that \mathbf{K} is *dense*, with non-zero off-diagonal elements. Such problems are also *globally data dependent*, in the sense that the calculation at each matrix element of the discretized system in the depends on all other elements. This density made numerical methods based on boundary integral equations prohibitively expensive prior to the discovery of so called ‘fast algorithms’, of which the FMM is the prototypical example. The naive computational complexity of storing a dense matrix, or calculating its matrix vector product is $O(N^2)$, and the complexity of finding its inverse is $O(N^3)$ with linear algebra techniques such as LU decomposition, where N is the number unknowns.




Figure 1.1: Given two boxes in \mathbb{R}^2 , \mathcal{B}_1 , \mathcal{B}_2 , which enclose corresponding degrees of freedom, off diagonal blocks in the linear system matrix $\mathbf{K}_{\mathcal{B}_1\mathcal{B}_2}$ and $\mathbf{K}_{\mathcal{B}_2\mathcal{B}_1}$ are considered low-rank for the FMM when separated by a distance at least equal to their diameter, this is also known as ‘strong admissibility’. Figure adapted from [45].

The critical insight behind the FMM, and other fast algorithms, is that one can compress physically distant interactions by utilising the rapid decay behaviour of the problem’s kernel. A compressed ‘low rank’ representation can be sought in this situation, displayed in figure (1.1) for \mathbb{R}^2 .

Using the FMM the best case the matrix vector product described by (1.1) and (1.3) can be computed in just $O(N)$ flops and stored with $O(N)$ memory. Fast algorithm based matrix inversion techniques display similarly optimal scaling in the best case. Given the wide applicability of boundary integral equations to natural sciences, from acoustics [67, 29] and electrostatics [66] to electromagnetics [14] fluid dynamics [51] and earth science [10]. Fast algorithms can be seen to have dramatically brought within reach large scale simulations of a wide class of scientific and engineering problems. The applications of FMMs aren’t restricted to BIEs, as (1.1) shares its form with the kernel summations often found in statistical applications, the FMM has found uses in and computational statistics [3], machine learning [35] and Kalman filtering [37]. The uniting feature of these applications is their global data dependency.

Recent decades have seen the development of numerous mathematical techniques for the computation of fast algorithms. However given their broad applicability across various fields of science and engineering, this has not been met with a commensurate development of black box open-source software solutions which are easy to use and deploy by non-experts. This is not to say that there is an absence of research software for the FMM [31, 69, 65, 39, 48], or fast matrix inversion [50, 44, 30]. However, the software landscape is heavily fragmented, codes often arising out of a software or mathematical investigation with infrequent maintenance or development post-publication. Few attempts have been made to re-use data structures, or application programming interfaces (APIs) between projects, and source code is often poorly documented leading to little to no interoperability between projects. Furthermore, as codes are often written in compiled languages such as Fortran [48] or C++ [39, 69, 65], there is a relatively high software engineering barrier entry for community contributions, further discouraging widespread adoption amongst non-specialist academics and industry practitioners. Additionally, significant domain specific expertise in numerical analysis is required by users to discern the subtle differences between fast algorithm implementations, or indeed to write one independently.

Computer hardware and architectures continue to advance concurrently with advances in numerical algorithms. Recently, the exascale benchmark (capable of 10^{18}

flops) was achieved by Oak Ridge National Labs' Frontier machine¹. With 9,472 AMD 64 core Trento nodes with a total of 606,208 compute cores, alongside 37,888 Radeon Instinct GPUs with a total of 8,335,360 cores, programming fast algorithms with their inbuilt global data dependency is challenging at a software level due to the communication bottlenecks imposed by the necessary all to all communications. Furthermore, the dense matrix operations required by fast algorithms require delicate tuning to fully take advantage of memory hierarchies on each node. Currently there exist very few open-source fast algorithm implementations that are capable of being deployed on parallel machines [39, 69], or take advantage of a heterogeneous CPU/GPU environments [69]. In fact for fast inverses there doesn't yet exist an open-source parallel implementation. Furthermore, developers must using existing codes must employ careful consideration in order to successfully compile the software in each new hardware environment they encounter, from desktop workstations to supercomputing clusters.

Resultantly, researchers who may want to write application code that takes advantage of fast algorithms as a black box without the necessary software or numerical analysis expertise to implement their own have few choices, and fewer still in a distributed computing setting. Identifying this as a significant barrier to entry for the adoption of fast algorithms in the wider community, we propose a new unified framework for fast algorithms, beginning with an implementation of a parallel FMM, which we introduce in the following section, designed for modern large scale supercomputing clusters. We emphasise our focus on ergonomic and malleable code, such that our code is easy to edit and deploy on a multitude of architectures while still achieving good scaling. With a key target application being the simulation of exascale boundary integral problems for electromagnetics, specified by Maxwell's equations.

1.2 From Analytic to Algebraic Hierarchical Fast Multipole Methods

The FMM as originally presented has since been extended into a broad class of algorithms with differing implications for practical implementations. We consider the problem in its most generic form by returning to the matrix vector product (1.1). We consider a non-oscillatory problem with a Laplace kernel from electrostatics to introduce the ideas behind the FMM, as in the original formulation [26]. Given an N -body evaluation of electrostatic potentials, we let $\{x_i\}_{i=1}^N \in \mathbb{R}^d$ denote the set of locations of charges of strength q_i , where $d = 2$ or $d = 3$. Our task is then to evaluate potentials, ϕ_j for $i = 1, 2, \dots, N$. We can without loss of generality take the value of $K(x, x) = 0$. Denoting our square domain containing all points with Ω , we seek a matrix vector product of the form,

$$\phi = \mathbf{K}\mathbf{q} \tag{1.4}$$

where $\phi \in \mathbb{C}^N$, $\mathbf{q} \in \mathbb{C}^N$ and $\mathbf{K} \in \mathbb{C}^{N \times N}$. The idea is to compress the kernel interactions defined by $K(x, y)$ when x and y are distant. Consider the situation in figure (1.3) where we choose \mathbb{R}^2 for simplicity. Here we seek to evaluate the potential induced by the source particles, $\{y_j\}_{j=1}^M$, in Ω_s at the target particles, $\{x_i\}_{i=1}^L$ in Ω_t .

¹<https://www.olcf.ornl.gov/frontier/>

$$\phi_i = \sum_{j=1}^L K(x_i, y_j) q_j, \quad i = 1, 2, \dots, M \quad (1.5)$$

As the sources and targets are physically distant, we can apply a low-rank approximation for the kernel as a sum of tensor products,

$$K(x, y) \approx \sum_{p=0}^{P-1} B_p(x) C_p(y), \quad \text{when } x \in \Omega_t, y \in \Omega_s \quad (1.6)$$

where P is called the ‘expansion order’, or ‘interaction rank’. We introduce the index sets I_s and I_t which list the points inside Ω_s and Ω_t respectively, and consider a generic approximation by tensor products where,

$$\hat{q}_p = \sum_{j \in I_s} C_p(x_j) q_j, \quad p = 0, 1, 2, \dots, P-1 \quad (1.7)$$

this is valid as K is smooth in the far field. Using this, we evaluate the approximation of the potential at the targets as,

$$\phi_i \approx \sum_{p=0}^{P-1} B_p(x_i) \hat{q}_p \quad (1.8)$$

In doing so we see that we accelerate (1.5) from $O(ML)$ to $O(P(M+L))$. As long as we choose $P \ll M$ and $P \ll L$, we will recover an accelerated matrix vector product. The power of the FMM, and similar fast algorithms, is that we can recover the potential in Ω_t with high accuracy even when P is small. We deliberately haven’t stated how we calculate B_p or C_p . In Greengard and Rokhlin’s FMM these took the form of analytical multipole and local expansions of the kernel function [26].

To demonstrate this we derive an expansion in the \mathbb{R}^2 case, taking c_s and c_t as the centres of Ω_s and Ω_t respectively,

$$\begin{aligned} K(x, y) &= \log(x - y) = \log((x - c_s) - (y - c_s)) \\ &= \log(x - c_s) + \log\left(1 - \frac{y - c_s}{x - c_s}\right) \\ &= \log(x - c_s) - \sum_{p=1}^{\infty} \frac{1}{p} \frac{(y - c_s)^p}{(x - c_s)^p} \end{aligned} \quad (1.9)$$

where the series converges for $|y - c_s| < |x - c_s|$. We note (1.9) is exactly of the form required with $C_p(y) = -\frac{1}{p}(y - c_s)^p$ and $B_p(x) = (x - c_s)^{-p}$. We define a ‘multipole expansion’ of the charges in Ω_s as a vector $\hat{\mathbf{q}}^s = \{\hat{q}_p^s\}_{p=0}^{P-1}$,

$$\begin{cases} \hat{q}_0^s = \sum_{j \in I_s} q_j \\ \hat{q}_p^s = \sum_{j \in I_s} -\frac{1}{p} (x_j - c_s)^p q_j, \quad p = 1, 2, 3, \dots, P-1 \end{cases} \quad (1.10)$$

The multipole expansion is a representation of the charges in Ω_s and can be truncated to any required precision. We can use the multipole expansion in place of a direct calculation with the particles in Ω_s . As the potential in Ω_t can be written as,

$$\phi(x) = \sum_{j \in I_s} K(x, y) q_j = \log(x - c_s) \hat{q}_0^s + \sum_{p=1}^{\infty} \frac{1}{(x - c_s)^p} \hat{q}_p^s \quad (1.11)$$

Greengard and Rokhlin also define a local expansion centered on Ω_t , that represents the potential due to the sources in Ω_s .

$$\phi(x) = \sum_{p=1}^{\infty} (x - c_t)^p \hat{\phi}_p^t \quad (1.12)$$

with a simple computation to derive the local expansion coefficients $\{\hat{\phi}_p^t\}_{p=0}^{\infty}$ from $\{\hat{q}_p^s\}_{p=0}^{P-1}$ (see app. A.1).

For our purposes it's useful to write the multipole expansion in linear algebraic terms as a linear map between vectors,

$$\hat{\mathbf{q}}^s = \mathbf{T}_s^{P2M} \mathbf{q}(I_s) \quad (1.13)$$

where \mathbf{T}_s^{P2M} is a $P \times N_s$ matrix, analogously for the local expansion coefficients we can write,

$$\hat{\phi}^t = \mathbf{T}_{t,s}^{M2L} \hat{\mathbf{q}}^s \quad (1.14)$$

where $\mathbf{T}_{t,s}^{M2L}$ is a $P \times P$ matrix, and the calculation of the final potentials as,

$$\phi^t = \mathbf{T}_t^{L2P} \hat{\phi}^t \quad (1.15)$$

where \mathbf{T}_t^{L2P} is a $N_t \times P$ matrix. Here we denote each *translation* operator, \mathbf{T}^{X2Y} , with a label read as ‘X to Y’ where L stands for local, M for multipole and P for particle. Written in this form, we observe that one could use a different method to approximate the translation operators than explicit kernel expansions to recover our approach’s algorithmic complexity, and this is indeed the main difference between different implementations of the FMM.

We have described how to obtain linear complexity when considering two isolated nodes, however in order to recover this for interactions between *all particles* with we rely on a hierarchical partitioning of Ω using a data structure from computer science called a *quadtree* in \mathbb{R}^2 or an *octree* in \mathbb{R}^3 . The defining feature of these data structures is a recursive partition of a bounding box drawn over the region of interest (see fig. 1.2). This ‘root node’ is subdivided into four equal parts in \mathbb{R}^2 and eight equal parts in \mathbb{R}^3 . These ‘child nodes’ turn are recursively subdivided until a user defined threshold is reached based on the maximum number of points per leaf node. These trees can be ‘adaptive’ by allowing for non-uniform node sizes, and ‘balanced’ to enforce a maximum size constraint between adjacent nodes [60].



Figure 1.2: An adaptive octree for random point data placed on the surface of a ‘wiggly torus’ test geometry. The user defines the level of recursion via a threshold for the maximum number of particles in a given node.



Figure 1.3: Two well separated clusters Ω_t and Ω_s where we can apply a low-rank approximation.

In addition to the \mathbb{T}^{P2M} , \mathbb{T}^{M2L} and \mathbb{T}^{L2P} the FMM also require operators that can translate the expansion centre of a multipole or local expansion, \mathbb{T}^{L2L} , \mathbb{T}^{M2M} , an operator that can form a local expansion from a set of points \mathbb{T}^{P2L} , and apply a multipole approximation to a set of points, \mathbb{T}^{M2P} , finally we need define a $P2P$ operator as short hand for direct kernel evaluations. Algorithm (1) provides a brief sketch of the full FMM algorithm.

In its original analytical form the applicability of the FMM is limited by the requirement for an explicit multipole and local expansions, as well as a restriction to matrix vector products. This lead to the development of algebraic variants that operate on the matrix represented by (1.4) directly, without the need for geometrical considerations via an octree. Examples include, the \mathcal{H} matrix [28], \mathcal{H}^2 matrix [8], hierarchically semi-separable (HSS) [11], and hierarchically off-diagonal low-rank (HODLR) [2] matrices.

These methods all represent the system matrix in (1.4) using a stored hierarchical matrix factorisation. Once this factorisation is computed, it can be used again - allowing for simple extensions to matrix-matrix products. Furthermore, an explicit

Algorithm 1 Adaptive Fast Multipole Method: FMM literature distinguishes between types of relationships between neighbouring nodes with the concept of *interaction lists*. There are four such lists for a given node B , called V , U , W and X . For a leaf node B , the U list contains B itself and leaf nodes adjacent to B . and the W list consists of the descendants of B 's neighbours whose parents are adjacent to B . For non-leaf nodes, the V list is the set of children of the neighbours of the parent of B which are not adjacent to B , and the X list consists of all nodes A such that B is in their W lists. The non-adaptive algorithm is similar, however the W and X lists are empty.

N is the total number of points

s is the maximum number of points in a leaf node.

Step 1: Tree construction

for each node B in *preorder* traversal of tree **do**

subdivide B if it contains more than s points.

end for

for each node B in *preorder* traversal of tree **do**

construct *interaction lists*, U , V , X , W

end for

Step 2: Upward Pass

for each leaf node B in *postorder* traversal of the tree **do**

P2M: compute multipole expansion for the particles they contain.

end for

for each non leaf node B in *postorder* traversal of the tree **do**

M2M: form a multipole expansion by translating the expansion centre of its children to its centre and summing their multipole expansion coefficients.

end for

Step 3: Downward Pass

for each non-root node B in *preorder* traversal of the tree **do**

M2L: translate multipole expansions of nodes in B 's V list to a local expansion at B .

P2L: translate the charges of particles in B 's X to the local expansion at B .

L2L: translate B 's local expansion to its children by translating its expansion centre to the centre of its children, and assigning the same coefficients.

end for

for each leaf node B in *preorder* traversal of the tree **do**

P2P: directly compute the local interactions using the kernel between the particles in B and its U list.

L2P: translate local expansions for nodes in B 's W list to the particles in B .

M2P: translate the multipole expansions for nodes in B 's W list to the particles in B .

end for




Figure 1.4: Matrix containing common algebraic alternatives to the FMM, adapted from [1].

matrix forms also allows for optimal matrix inversion algorithms.

Consider an index set I , corresponding to the indices of all points $\{x_i\}_{i=1}^N$ in a given discretisation. The general approach of these methods is to partition I in such a way that we can exploit the low-rank interactions between distantly separated clusters of particles. Initially, an n -ary ‘cluster tree’ \mathcal{T}_I , with a set of nodes T_I , is formed such that

1. $T_I \subseteq \mathcal{P}(I) \setminus \{\emptyset\}$, meaning each node of \mathcal{T}_I is a subset of the index set I , here $\mathcal{P}(I)$ denotes the power set of I .
2. I is the root of \mathcal{T}_I
3. The number of indices in a leaf node $\tau \in T_I$ is such that $|\tau| \leq C_{\text{leaf}}$ where C_{leaf} is a small constant.
4. Non leaf nodes τ have n child nodes $\{\tau_i\}_i^n$, and is formed of their disjoint union $\tau = \cup_{i=1}^n \tau_i$

Cluster trees may in general be n -ary, however common implementations are as binary trees. Using a cluster tree, one forms a ‘block tree’, $\mathcal{T}_{I \times I}$. Each node, or ‘block’, of a block tree, $N(\tau \times \sigma)$, corresponds to the clusters represented by indices $\tau \subseteq I$ and $\sigma \subseteq I$ respectively. The \mathcal{H} and \mathcal{H}^2 representations are ‘strongly admissible’, meaning that well separated blocks as in figure (1.1) can be compressed. In contrast, the HSS and HODLR approaches are ‘weakly admissible’, and also consider adjacent clusters to be compressible. Admissibility is calculated by forming a bounding box around clusters, and checking their separation along each dimension [7]. Furthermore, as in the FMM which creates parent multipole expansions for a given node from that of its children, and vice versa for local expansions, an analogous approach is taken by \mathcal{H}^2 and HSS matrices, referred to as ‘nested bases’. These different approaches are succinctly visualised in figure (1.4).

Expressed in this way the matrix implicit in the FMM (alg. (1)), can be seen to be a member of the \mathcal{H}^2 class of matrices. Therefore the algorithms developed for algebraic approaches for matrix vector products similarly recover optimal $O(N)$ scaling in the best case, with the additional benefit of easy extension to matrix matrix products, and matrix inversion in optimal complexity [7].

From a computational perspective, the trade-off between these analytical and algebraic approaches for the FMM are best expressed in terms of the ratio of compute (flops) to memory (bytes), or ‘arithmetic intensity’. The analytical FMM has a high arithmetic intensity, due to its matrix free nature. However, using explicit \mathcal{H}^2 representations requires one to compute and store the full hierarchical matrix in

memory, resulting in increased memory movement costs, both vertically on a single node, and horizontally in a distributed memory setting.

A third ‘semi-analytical’ method, combining the best of the purely analytical and algebraic methods are known as ‘black box’, or ‘kernel independent’, FMMs [36, 18, 40]. These methods mimic the algorithmic structure of the analytical FMM, with its matrix free nature, however the low rank representations they construct are done so independently of the kernel, hence the name ‘black box’. We contrast two common black box methods in box (1.5), the underlying difference being in their representation of translation operators. Black box methods are preferable from a software standpoint for three reasons:

1. Generic software interfaces can be built for a variety of kernels, unlike for analytical FMMs, as demonstrated by [65, 31].
2. They have reduced storage requirements in comparison to purely algebraic approaches.
3. The majority of computations are simple matrix vector products that are readily expressed with BLAS L2 operations, and therefore maximise cache efficiency

We emphasise that it is possible to extend the analytical FMM to more general kernels, such as the modified Laplacian [27], Stokes [20] and Navier [21], however software implementations are cumbersome, requiring the optimised calculation of special functions such as spherical harmonics, and are difficult to generalise for different problem settings. We have preferentially implemented the Kernel-Independent FMM of Ying et. al [36] in our previous software [31] as it has a particularly simple mathematical structure, and extends to a variety of kernels for elliptic partial differential equations of interest. Specifically, the Laplace, Stokes and Navier kernels alongside their modified variants [36], as well as the Helmholtz kernel in the low-frequency setting [65].

In terms of algorithm optimisation significant efforts have gone towards optimising T^{M2L} [42, 18, 36], which is the most time consuming translation operator. It is computed for each non-leaf node between level one and the leaf level of a tree, it therefore grows as $O(N)$. In \mathbb{R}^3 , the V list for a node, containing its parent’s neighbors’ children that are non-adjacent to the node, can be up to $|V| = 6^3 - 3^3 = 189$. However if a kernel exhibits translational invariance, such that

$$K(x, y) = K(x - y)$$

which is the case for the kernels mentioned above, we recognise that there are just $7^3 - 3^3 = 316$ unique T^{M2L} operators per level. However, Messner et. al [42] show that even these 316 operators are permutations of a subset of only 16 unique operators. Permutations corresponding to reflections across various planes (for example $x = 0$, $y = 0$, $z = 0$, etc). We can therefore write a block matrix of *all* unique T^{M2L} for a given level,

$$T^{M2L} = [T_1^{M2L} | T_2^{M2L} | \dots | T_{16}^{M2L}] \quad (1.16)$$

where T_i^{M2L} corresponds to the i^{th} *transfer vector* describing a unique T^{M2L} . This block matrix is efficiently compressed using a truncated SVD, which is known

to produce an optimal low rank approximation of a rank N matrix, where the new rank, r , is chosen such that $r \ll N$. This is a common optimisation used in both black box [36, 18] and analytical FMMs [25]. By precomputing and storing a compressed T^{M2L} for each level, the T^{M2L} operator for each node can be formed as a single BLAS level 2 operation at runtime by looking up the relevant compressed operators corresponding to its V list. The T^{M2M} , T^{L2L} can be similarly stored on a level by level basis, reducing the pre-computation time for operators to $O(d)$, where d is the depth of an quad/octree. For homogenous kernels, for which,

$$K(\alpha x, \alpha y) = \alpha^m K(x, y)$$

pre-computations can be performed for a single level, and scaled to different levels, reducing pre-computation complexity to $O(1)$. Such optimisations further separate runtime performance between implementations of algebraic and semi-analytical methods, as the transfer operators T^{M2L} correspond to duplicate blocks in an algebraic \mathcal{H}^2 which are redundantly computed and stored in implementations [24].


A significant challenge in practical FMM implementations is handling the global data dependency encapsulated in (1.5). Historical single-core architectures have always suffered from the Von Neumann bottleneck, whereby data is loaded from main memory slower than the processor is able to absorb it. However, with the breakdown of Dennard scaling² from around 2006 and the subsequent multicore revolution the disparity between data movement costs and computation have been even further exacerbated. On modern hardware accelerators are able to achieve $O(1e12)$ flop rates in double precision, but are bandwidth limited to $O(1e11)$ bytes/second, meaning that an order of magnitude more operations have to be performed for every byte retrieved from memory, to remain efficient. Dongarra et. al [16], go as far as to suggest that the dominant role of communication costs on emerging hardware, especially in the context of exascale architectures, mean that complexity analysis of distributed algorithms based on flop counts are redundant entirely. Indeed, in a distributed memory setting global reductions are often the largest performance limiting factor [16]. The data dependency of the FMM is expressed in the quad/octree, with significant global reductions required for translation operators. It is therefore critical to implement highly optimised algorithms for tree construction in a distributed setting, however we defer a discussion on this until section 2.4.

We conclude by mentioning that the FMM and its variants are not the only technique available to accelerate (1.1), for problems in which only uniform resolution is required the FFT, which has corresponding distributed memory implementations [22]. However, we are commonly interested in multiscale problems in which neighbouring nodes can be of differing sizes. In this setting, multigrid methods have shown slower convergence in comparison to the FMM [68, 23]. Furthermore, a multigrid approach does not allow for a re-use of FMM data structures (sec. 3.3) for other fast algorithms we are interested in implementing, and aim to present as a unified framework with maximum code re-use. We observe that the choice between analytical, semi-analytical and algebraic FMM variants may have a significant impact on the ultimate implementation, and performance, this is something we explicitly measure in section 2.1. We conclude that the semi-analytical Kernel-Independent FMM of Ying et. al [36], offers a good compromise between ease of software design, mathematical simplicity, and problem generality.

²Dennard Scaling is the insight that the power density of transistors appeared to remain constant as they grew smaller.

Kernel-Independent FMM (KIFMM)


The KIFMM approximates the multipole expansion for a given leaf node containing particles $\{y_j\}_j^N$ with charges $\{q_j\}_j^N$ by evaluating their potential directly (1.4) at corresponding evenly spaced points on a ‘check surface’, displayed in blue.



This is matched to N_e ‘equivalent charges’, q^e , placed evenly on a the (red) equivalent surface at $\{x_i\}_i^{N_e}$. These surfaces are taken to be circles in \mathbb{R}^2 and cubes in \mathbb{R}^3 . They perform a Tikhonov-regularised least squares fit to calculate q^e , which plays the role of the *multipole expansion*.

$$q^e = (\alpha I + K^*K)^{-1}Kq^{\text{node}}$$

where q^{node} are the charges in the leaf node. Using a similar framework, involving equivalent and check surfaces, the T^{M2M} , T^{M2L} , T^{L2L} and T^{L2P} operators are also calculated with Tikhonov-regularised least squares fitting, and are sketched below.



The KIFMM is designed for non-oscillatory kernels, such as the Laplace kernel of our didactic example, though in practice works well with low-frequency oscillatory problems, with extensions to high frequency problems [17].

The least-squares fit can be pre-computed for each given geometry, and re-used. Additionally, as the KIFMM decomposes into a series of matrix vector products, it is well suited to software implementations.

Black-Box FMM (bbFMM)

An n point interpolation scheme for the kernel is constructed sequentially over each variable,

$$K(x, y) \approx \sum_{l=1}^n K(x_l, y)w_l(x)$$

$$K(x, y) \approx \sum_{l=1}^n \sum_{m=1}^n K(x_l, x_y)w_l(x)w_m(y)$$

with coefficients $w_l(x)$ and $w_m(y)$. The bbFMM uses first kind Chebyshev polynomials, T_n to interpolate the kernel, defined as,

$T_n(x) = \cos(n\theta)$, where $x = \cos(\theta)$ over the closed interval $x \in [-1, 1]$. The roots, $\{\bar{x}_m\}$, known as Chebyshev nodes are defined as,

$$\bar{x}_m = \cos(\theta_m) = \cos\left(\frac{(2m-1)\pi}{2n}\right)$$

This is a well known, stable, uniformly convergent, interpolation scheme, that doesn’t suffer from Runge’s phenomenon [18]. An n point Chebyshev approximation to a given function, $g(x)$ with $p_{n-1}(x)$, can be written as,

$$p_{n-1}(x) = \sum_{k=1}^{n-1} c_k T_k(x)$$

where, $c_k = \begin{cases} \frac{2}{n} \sum_{l=1}^n g(\bar{x}_l) T_k(\bar{x}_l), & \text{if } k > 0 \\ \frac{1}{n} \sum_{l=1}^n g(\bar{x}_l), & \text{if } k = 0 \end{cases}$

we recognise,

$$p_{n-1}(x) = \sum_{l=1}^n g(\bar{x}_l) S_n(\bar{x}_l, x)$$

$$\text{with, } S_n(x, y) = \frac{1}{n} + \frac{2}{n} \sum_{k=1}^{n-1} T_k(x) T_k(y)$$

When applied to our generic interpolation for the kernel,

$$K(x, y) \approx \sum_{l=1}^n \sum_{m=1}^n K(\bar{x}_l, \bar{y}_m) S_n(\bar{x}_l, x) S_n(\bar{y}_m, y)$$

which can be substituted into (1.5), recovering linear complexity as long as the number of Chebyshev nodes is small, $n \ll N$.

Figure 1.5: The formulations of the kernel-independent FMM methods, the KIFMM [36] and the bbFMM [18].

Designing Software for Fast Algorithms

2.1 The Software Landscape

Numerous different groups have developed overlapping software for fast algorithms. For FMMs, leading analytical implementations include the single-node multithreaded, ‘FMM3D’ and ‘FMM2D’ [48, 49], and the MPI accelerated ‘ExaFMM’ [69] packages. The ExaFMM project have also developed a single-node multithreaded kernel-independent FMM, ‘ExaFMM-T’ [65]. For algebraic methods, the landscape is sparser, with few parallel implementations. The standouts being the MPI accelerated ‘STRUMPACK’ [24], for HODLR and HSS matrices, as well as the multithreaded single-node ‘FLAM’ [30], which supports algorithms for \mathcal{H}^2 matrices. The majority of codes, except FLAM, are written in either Fortran or C++, with a few supporting interfaces to higher level languages such as Python and Matlab [65, 48].

We compare these softwares for the calculation of (1.5) with N randomly distributed particles placed in a unit cube, $[0, 1]^3$, with uniform charge, $q_i = 1$, in figure (2.1)¹. Experiments are taken on a single-node AMD Ryzen Threadripper 3970X 32 core Processor, with 250GB of memory. To avoid thread oversubscription in STRUMPACK and ExaFMM, both designed for multi-node systems, the maximum number of OpenMP threads is restricted to one. Each FMM software is run with an expansion order $P = 10$, and the algebraic software’s parameters are adjusted to match this level of accuracy.

In comparing these softwares we aimed to emulate the workflow of a typical user evaluating the differences between software packages, and therefore restricted our study to runs which converged in less than 30 minutes, $\approx 2 \times 10^3$ s. Figure (2.1a) plots the time to compute (1.5), including the time to create all required data structures. For the algebraic softwares, FLAM and STRUMPACK, this includes the time to factorise and store an explicit matrix representation of (1.5). Strictly speaking, STRUMPACK’s HSS and HODLR matrices are not applicable to (1.5), as it is a strongly admissible problem. We see the effects of this in super-quadratic factorisation times for increasing N . Additionally, FLAM, written in Matlab, is not designed for high performance. However, these remain the only well supported, publicly available, parallel implementations of algebraic fast algorithms and we include them for comparison. As expected, the factorization scales super-linearly for both softwares, which lies in contrast with their relatively fast matrix vector products once factorised (fig. (2.1c)).

We observe that currently available open source algebraic software is impractical for computing even moderately sized matrix vector products, i.e. $N \geq O(10^5)$.

¹Our experiments are available at https://github.com/skailasa/phd-thesis/code/ch_2/

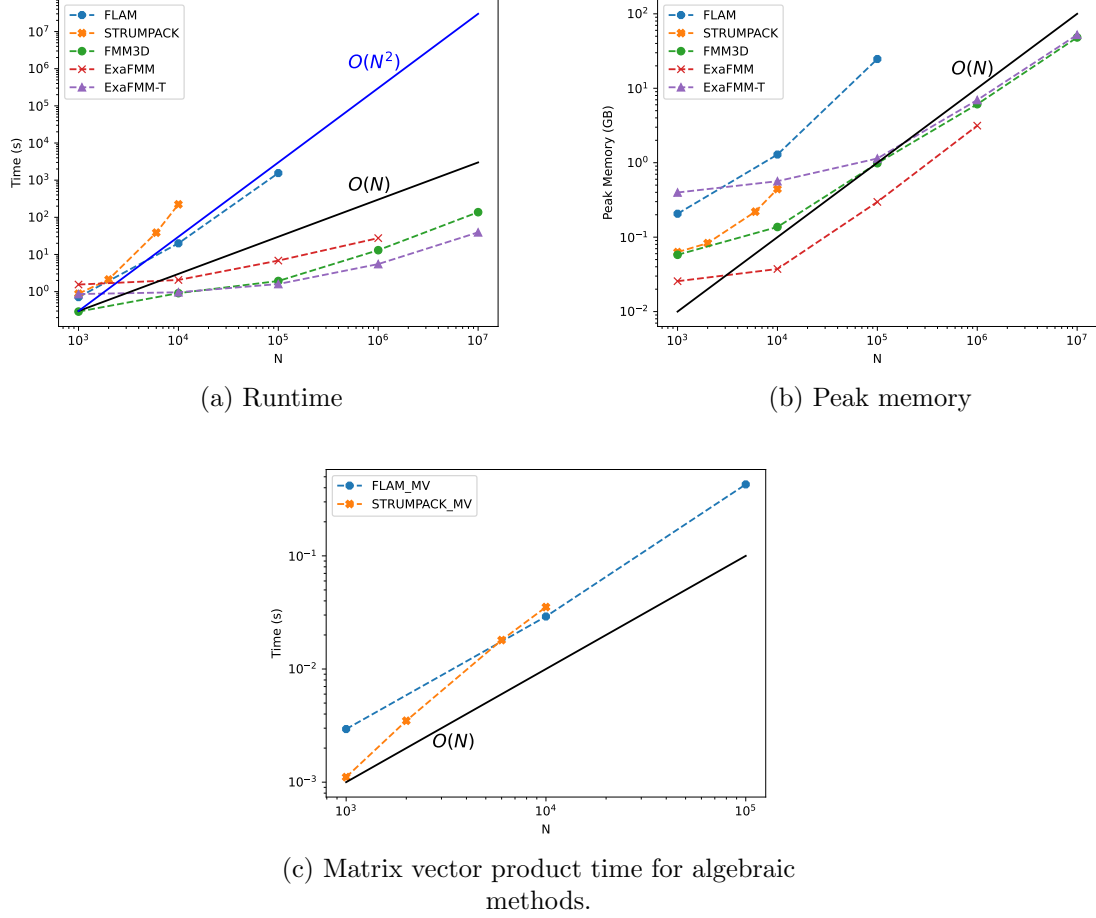


Figure 2.1: Comparison of parallel software for fast algorithms, used to calculate (1.5). Simulations were compared for convergence times $\leq 2 \times 10^3$ s that could fit into memory.

With runtimes times for factorisation exceeding our imposed limit of 2×10^3 s. One might argue that their expensive factorization cost is worth it if one proceeds to compute (1.5) multiple times, or if one wants to quickly find an inverse. However, the FMM softwares tested beat both algebraic softwares by several orders of magnitude, making it much more tenable to apply them multiple times if required, and to find matrix inverses via an iterative Krylov type methods.

As expected from theory, the analytical FMMs, ExaFMM and FMM3D, consume less memory than the kernel-independent ExaFMM-T (fig. (2.1b)). However, for larger problem sizes this difference is marginal between FMM3D and ExaFMM-T. For the largest problem size tested, $N = 10^7$, ExaFMM-T is roughly 3.5 times faster than FMM3D with a similar memory footprint. Counter intuitively, the analytical ExaFMM fails to converge in our experiments for $N = 10^7$, as its memory requirements exceed the available 250 GB.

Theoretical considerations aside, each software's performance is practically determined by the computational and algorithmic performance optimisations they take. ExaFMM-T for example uses optimized algorithms for the calculation of inverse square roots [39], as well as using SIMD vectorisation for kernel evaluations [65]. Similarly by stacking T^{M^2L} as in (1.16), they optimise cache-reuse. Similar SIMD vectorisations are implemented by FMM3D and ExaFMM, however their remaining optimisations are presented opaquely, making it difficult for users to evaluate their

relative merits.

Usability, defined in terms of quality of API documentation, available descriptions of mathematical and software optimisations, quality of software engineering in terms of self-documenting and well commented code, as well as ease of installation, is as critical to the dissemination of scientific software as raw performance. It's unlikely that a user will incorporate a performant piece of software into their work if it's not usable. In this respect FMM3D and ExaFMM-T stand out, with well documented APIs, code examples, as well as wrappers to call functionality from higher level languages. STRUMPACK, though relatively well documented, does not support wrappers to higher-level languages. Written in C++, its documentation is largely a description of its object hierarchy, with few examples documenting real world use cases. ExaFMM on the other hand is poorly documented, with few tests or comments, and no wrappers to high-level languages. Most softwares, except STRUMPACK which is distributed via the SPACK package manager, have traditional source builds based on CMake and Make. Local installation can therefore be challenging when building on non-traditional HPC operating systems such as Windows and MacOS.

From the experiments in figure (2.1) it's clear that software choice can have dramatic impact on runtime and memory scaling regardless of algorithmic choice. For the Laplace problem of (1.5) all softwares can be used more or less out of the box, however oscillatory Helmholtz kernels are only natively supported by a few [69, 65, 49, 48]. The most complete functionality, with the addition of Stokes and Maxwell kernels, are only supported by FMM2D and FMM3D. From a non-expert user's perspective, merely understanding the theory behind each algorithm isn't enough to expect a guaranteed performance, as demonstrated by the poor scaling of algebraic softwares, as well as the different memory requirements of the analytical FMMs. In light of this, users would likely be tempted to pick a 'median' option, which guarantees good performance, scalability, easy installation, and usability via a high-level language wrapper and support for multiple problems. This makes FMM2D and FMM3D the standout options, however as they're developed as a high-performance Fortran libraries, they are relatively unmalleable, with a steep learning curve for potential contributors and limited to single-node systems.

In conclusion, we identify a distinct lack of software for fast algorithms that fulfills our usability criteria and can also scale to multi-node systems. Presently available software is not composable, with differing APIs and re-implementations of high-performance kernels as well as data structures such as quad and octrees. Algebraic methods lack an optimised software for strongly admissible problems, with FLAM designed as a research code for prototyping purposes. Wrappers for high-level languages are not uniformly available, and builds can be challenging in non-Linux environments.

We aim to solve this with our proposed software infrastructure (fig (2.2)). By offering a set of composable libraries with a unified API for FMMs and algebraic methods for constructing matrix inverses, we aim to maximally re-use highly-optimised code, such as for SIMD vectorised kernel evaluations, and quad and octrees that can scale across multiple processors. We propose Rust as the language for our this infrastructure. We initially explored high-performance code generation tools for Python as a bridge between an ergonomic language for researchers and developers with high-performance. This is documented in section 2.2 via our attempt to build a Python FMM. Due to the constraints of development in high-level languages, and not wishing to use C/C++ or Fortran, we pivoted to Rust. Rust is a young low-level

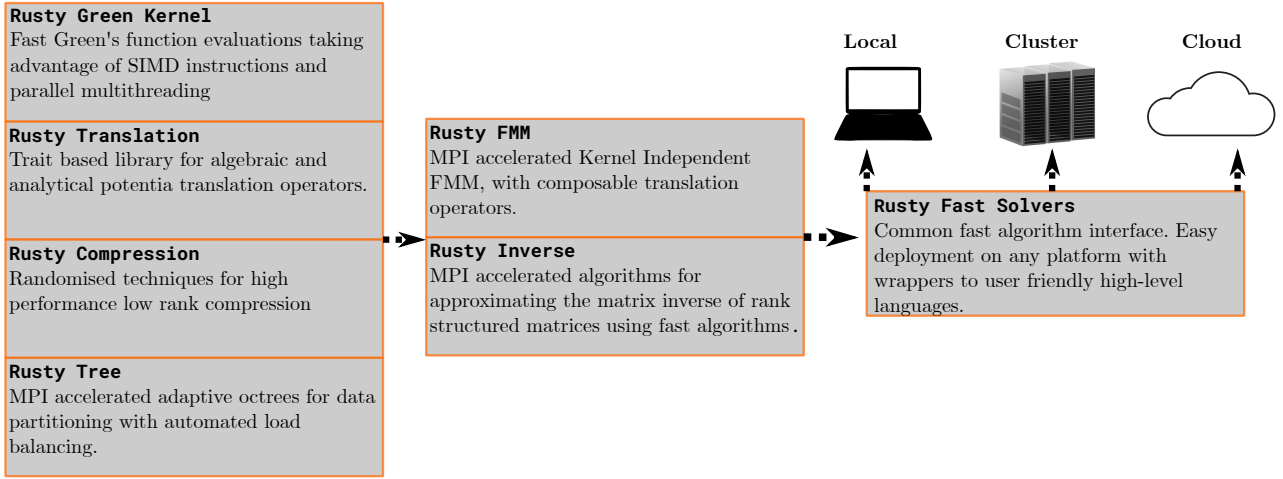


Figure 2.2: Hierarchy of our proposed infrastructure, with key functionality separated into individual libraries we maximise code re-use across projects. Once completed using Rust’s code generation infrastructure, users will be able to deploy our software from desktop workstations, to supercomputing clusters and expect scalable high-performance code, while being able to write application code in Python.

systems programming language, with significant industry support, which has many features that improve the usability of scientific software. Most importantly its build system, Cargo, is specified by a simple TOML file, making it easy to deploy Rust codes on most common platforms without the need for complex build CMake style build scripts. This will make it easy for users to develop software locally, before deploying to a super computing cluster, while still being able to expect good performance. Rust’s syntax inherits from both functional and imperative paradigms, and its trait system wholly replaces object orientation, making the description of shared behaviour, and data oriented programming, significantly more readable. However, importantly, it is also easy to develop wrappers from Rust to Python using its C foreign function interface [41], allowing us to expose our infrastructure to novice programmers. We document the prominent features of Rust that make it useful for scientific computing in section 2.3, before exploring the performance of a completed Rust based library for MPI accelerated octrees in section 2.4.

2.2 Case Study: PyExaFMM, a Python Fast Multipole Method

This section is adapted from a paper recently submitted to Computing in Science and Engineering [31]. Python², is designed for memory safety and programmer productivity. Its simplicity allows Computational Scientists to spend more time exploring their science, and less time being confused by software quirks, memory errors, and the nightmare of incompatible dependencies, which conspire to drain productivity in lower level languages.

Python fulfills many of the key usability criteria for scientific software. Cross platform builds are trivial with open source build systems such as Conda, and its simple syntax and large scientific computing ecosystem of numerical libraries allows

²We use ‘Python’ to refer to CPython, the popular C language implementation of Python, which is dominant in computational science.

for rapid dissemination amongst the wider community.

With this in mind we sought to assess the suitability of Python as a base language for fast algorithm software. With stable MPI bindings for multi-node development [13], Python’s major pitfall is its restriction to run in a single thread due a software construction called the ‘Global Interpreter Lock’ (GIL). Libraries for high performance computational science have traditionally bypassed the issue of the GIL by using Python’s C interface to call extensions built in C or other compiled languages which can be multithreaded or compiled to target special hardware features, thus enabling hierarchical parallelism.

Recently ‘just in time’ (JIT), compilers have emerged as a technique for compiling high-level languages to fast machine code. The idea being that a user is able to rapidly iterate on their algorithm, with the compiler taking responsibility to deliver performance. The Numba compiler For Python was written to targets and optimise code written with NumPy’s n -dimensional array, or ‘ndarray’, data structure, which are homogenously typed containers stored contiguously in memory [33]. Its power comes from the ability to generate multithreaded architecture optimised compiled code while *only writing Python*. The promise of Numba is the ability to develop applications with speed that can rival C++ or Fortran, while retaining the simplicity and productivity of working in Python. We tested Numba’s suitability by developing ‘PyExaFMM’, an implementation of the three-dimensional Kernel Independent FMM [31].

Numba

Numba is a compiler built with LLVM, a framework for building custom compilers, to target a subset of Python code that uses ndarrays. LLVM provides an API for generating machine code for different hardware architectures such as CPUs and GPUs and is also able to analyze code for hardware level optimisations such as auto vectorization, automatically applying them if they are available on the target hardware [34]. Additionally, LLVM generated code may be multithreaded - bypassing the issue of the GIL. Furthermore Numba is able to use the metadata provided by ndarrays describing their dimensionality, type and layout to generate code that takes advantage of the hierarchical caches available in modern CPUs [33]. Altogether, this allows code generated by Numba to run significantly faster than ordinary Python code, and often be competitive with code generated from compiled languages such as C++ or Fortran.

From a programmer perspective using Numba, at least naively, doesn’t involve a significant rewrite. Python functions are simply marked for compilation with a special decorator, see listings (2.2), (2.3) and (2.1) for example syntax. This encapsulates the appeal of Numba. The ability to generate high-performance code for different hardware targets from Python, and letting Numba worry about how to perform optimisations, would allow for significantly faster workflows than possible with a compiled language.

Figure (2.5) illustrates the program execution path when a Numba decorated function is called from the Python interpreter. We see that Numba doesn’t replace the Python interpreter. If a marked function is called at runtime, program execution is handed to Numba’s runtime which compiles the function on the fly with a type signature matching the input arguments. This is the origin of the term ‘just in time’ [JIT] to describe such compilers.

The Numba runtime interacts with the Python interpreter dynamically, and

control over program execution is passed back and forth between the two. There is a cost to this interaction from having to ‘unbox’ Python objects into types compatible with the compiled machine code, and ‘box’ the outputs of the compiled functions back into Python compatible objects. This process doesn’t involve re-allocating memory, however pointers to memory locations have to be converted and placed in a type compatible with either Numba compiled code or Python.

Numba’s Pitfalls

Since its first release Numba has been extended to compile most functionality from the NumPy library, as well as the majority of Python’s basic features and standard library modules³. However, if Numba isn’t able to find a suitable Numba type for each Python type in a decorated function, or it sees a Python feature it doesn’t yet support, it runs in ‘object mode’, handling all unknown quantities as generic Python objects. To ensure a seamless experience Numba does this without reporting it to the user, unless explicitly marked to run in ‘no Python’ mode (see listing (2.3) and (2.1) for example syntax). However, object mode is often no faster than vanilla Python, putting the burden on the programmer to understand when and where Numba works. As Numba influences the way Python is written it’s more akin to a programming framework rather than just a compiler.

An example of Numba’s framework-like behavior arises when implementing algorithms that share data, and have multiple logical steps as in listing (2.3). This listing shows three implementations of the same logic, a function that generates a random matrix $A \in \mathbb{R}^{100 \times 100}$ and multiplies it with itself, and also writes an input vector $v \in \mathbb{R}^{100}$ to a Numba dictionary. Data of this size is chosen to reflect the typical amount of computation in a task parallelized by PyExaFMM. The implementations algorithm 1 and algorithm 2 aren’t distinguished by the Numba compiler, and both pay a cost to call subroutines defined outside of their function body. However, algorithm 1 pays a small additional (un)boxing cost in order to manipulate a globally defined Numba compatible dictionary, in comparison to a locally defined one in algorithm 2. The *nested* function in algorithm 3 differs from the other two implementations, by defining its sub-routines within its function body, rather than calling externally defined functions. This is an example of an *inlining* optimisation, which is picked up by LLVM at compile time.

The runtimes of all three implementations are shown in table (2.1) for three contrasting problem sizes, and shows how inlining can have a significant impact on runtime for this algorithm⁴. This example is designed to illustrate how small changes to writing style can impact the performance of an algorithm written with Numba. We emphasize that the speedup obtained from inlining is dependent on the size of the data being operated on as well as the program logic. Other factors such as memory latency for large data, or the passing of execution control between Python and Numba, with small data may become more significant. Indeed the experiment with $A \in \mathbb{R}^{1000 \times 1000}$ and $v \in \mathbb{R}^{1000}$, we observe that inlining is still the dominant factor in performance difference and is even more prominent than with the smaller dataset. With $A \in \mathbb{R}^{1 \times 1}$ and $v \in \mathbb{R}^1$, we see that the instantiation of the result dictionary from within a Numba function is now a significant part of total

³A full list of supported features for the current release can be found at: <https://numba.pydata.org/numbadoc/dev/reference/pysupported.html>

⁴All experiments in this work were taken on an AMD Ryzen Threadripper 3970X 32-Core processor running Python 3.8.5 and Numba 0.53.0

runtime.

Algorithm	Matrix Dimension	Time (μ s)
1	$\mathbb{R}^{1 \times 1}$	1.74 ± 0.01
1	$\mathbb{R}^{100 \times 100}$	308 ± 1
1	$\mathbb{R}^{1000 \times 1000}$	27100 ± 200
2	$\mathbb{R}^{1 \times 1}$	2.94 ± 0.01
2	$\mathbb{R}^{100 \times 100}$	306 ± 1
2	$\mathbb{R}^{1000 \times 1000}$	27100 ± 200
3	$\mathbb{R}^{1 \times 1}$	2.61 ± 0.01
3	$\mathbb{R}^{100 \times 100}$	2.64 ± 0.07
3	$\mathbb{R}^{1000 \times 1000}$	2.84 ± 0.07

Table 2.1: Testing the effect of inlining and (un)boxing with dense matrix vector products in double precision using implementations from listing (2.3).

Nested functions have the tendency to grow long in performant Numba code, in order to minimize the number of interactions between Numba and Python. However this makes them more difficult to unit test. Indeed, performant Numba code can look decidedly un-Pythonic. Numba encourages fewer user created objects, performance critical sections written in terms of loops over simple array based data structures, and potentially long nested functions.

Furthermore, not every supported feature from Python behaves in a way an ordinary Python programmer would expect, which has an impact on program design. An example of this arises when using Python dictionaries, which are central to Python, but are only partially supported by Numba. As they are untyped, and can have any Python objects as members, they don't neatly fit into a Numba compatible type. Programmers can declare a Numba compatible 'typed dictionary', where the keys and values are constrained to Numba compatible types, and pass it to a Numba decorated function at low cost. However, using a Numba dictionary from the Python interpreter is *always slower* than an ordinary Python dictionary due to the (un)boxing cost when getting and setting any item.

Therefore, though Numba is advertised as an easy way of injecting performance into your program via a simple decorator, it can be seen to have its own learning curve. Achieving performance requires a programmer to be familiar with the internals of its implementation and potential discrepancies that arise when translating between Python and the LLVM generated code, which may lead to significant alterations in the design of algorithms and data structures.

Data Oriented Design

Data oriented design is about writing code that operates on data structures with simple memory layouts, such as arrays, in order to optimally take advantage of modern hardware features. The idea being that it is easier for programmers to optimise for cache locality and parallelization if the data structures are easier to map to the hardware. This contrasts with object oriented design, where although code is organized around data the focus is on user created types or 'objects', where memory layout is obfuscated by the potential complexity of an object, which can contain multiple attributes of different types. This makes it harder to write code that takes advantage of cache locality. Numba's focus on ndarrays strongly encourages

data oriented design principles, which are reflected in the design of PyExaFMM’s octrees as well as its API.

Octrees can either be ‘pointer based’ [65], or ‘linear’ [60] (see sec 2.4). A pointer based octree uses objects to represent each node, with fields for a unique id, contained particles, associated expansion coefficients, potentials, and pointers to their parent and sibling nodes. This makes searching for neighbors and siblings easy, as one has to just follow pointers⁵. The linear octree implemented by PyExaFMM represents nodes by a unique id stored in a 1D vector, all other data such as expansion coefficients, particle data, and calculated potentials, are also stored in 1D vectors. Data is looked up by creating indices to tie a node’s unique id to the associated data. This is an example of how Numba can affect design decisions, and make software more complex, despite the data structures being simpler.

Figure (2.3) illustrates PyExaFMM’s design. There is only a single Python object, ‘Fmm’, which acts as the API. It initializes ndarrays for expansion coefficients and calculated potentials, and its methods interface with Numba compiled functions for the FMM operators and their associated data manipulation functions. When sharing data we prefer nested functions, however we keep the operator implementations separate from each other, which allows us to unit test them individually. This means that we must have at least one interaction between Numba and the Python interpreter to call the near field, T^{P2M} , T^{L2P} , T^{M2P} and T^{P2L} operators, $d - 2$ interactions to call the T^{M2L} and T^{L2L} operators, and d interactions for the T^{M2M} operator where d is the depth of the octree. The most performant implementation would be a single Numba routine that interacts with Python just once, however this would sacrifice other principles of clean software engineering such as modularity, and unit testing. This structure has strong parallels with software designs that arise from traditional methods of achieving performance with Python by interfacing with a compiled language such as C or Fortran. The benefit of writing in Numba is that we can continue to write in Python. Though as seen above, performant Numba code may only be superficially Pythonic through its shared syntax.

Multithreading in Numba

Numba enables multithreading via a simple parallel for loop syntax (see listing (2.1)) reminiscent of OpenMP. Internally Numba can use either OpenMP or Intel TBB to generate multithreaded code. We choose OpenMP for PyExaFMM, as it’s more suited to functions in which each thread has an approximately similar workload. The threading library can be set via the `NUMBA_THREADING_LAYER` environment variable.

Numerical libraries such as NumPy and SciPy implement many of their mathematical operations using multithreaded compiled libraries internally, such as OpenBLAS or IntelMKL. Numba compiled versions of these operations retains this internal multithreading. This leads to *nested parallelism* when combined with a multithreaded region declared with Numba, as in listing (2.1). This is where a parallel region calls a function with another parallel region inside it. Threads created by the two functions are not coordinated by Numba, and this leads to *oversubscription*, where the number of active threads exceeds the CPU’s hardware capacity. Many threads are left idle while the CPU is forced to jump between threads operating on different data. This leads to broken cache locality, and hanging threads threads,

⁵‘Siblings’ are defined as nodes which share a parent, and ‘neighbours’ are defined as adjacent nodes which may not share a parent.

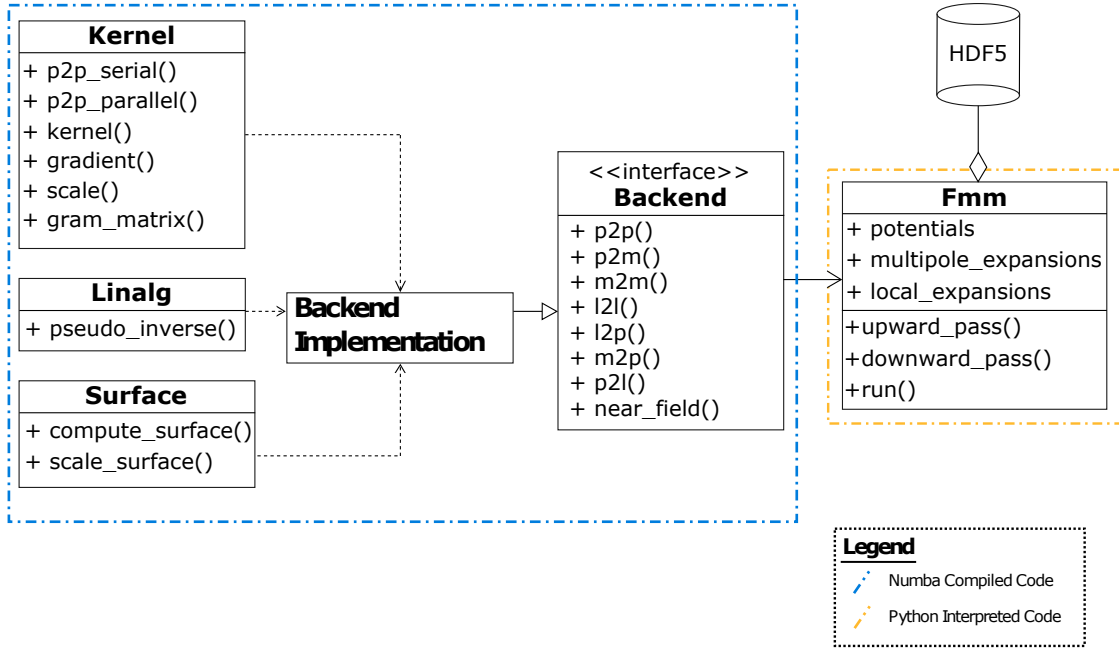


Figure 2.3: Simplified UML model of all PyExaFMM components. Trees and other precomputed quantities are stored in a HDF5 database. The ‘Fmm’ object acts as the user interface, all other components are modules consisting of methods operating on arrays, adapted from [31].

stuck waiting for others to finish [38]. We avoid this in PyExaFMM by explicitly setting NumPy operations to be single threaded, via the environment variable `OMP_NUM_THREADS=1`, before starting our program. This ensures that the only threads created are those explicitly declared using Numba.

Parallising the KIFMM

The T^{P2M} , T^{P2L} , T^{M2P} and T^{L2P} all rely on the P2P operator, as this computes (1.5) over their respective sources and targets, and are parallelized over their targets, the leaf nodes. For the T^{L2P} operator we encourage cache locality for the P2P step, and keep the data structures passed to Numba as simple as possible, by allocating 1D vectors for the source positions, target positions and the source expansion coefficients, such that all the data required to apply an operator to single target node is adjacent in memory. By storing a vector of ‘index pointers’, that bookend the data corresponding to each target in these 1D vectors, we can form parallel for loops over each target to compute the P2P that encourages cache-locality in the CPU. In order to do this, we have to first iterate through the target nodes, and lookup the associated data to fill the cache local vectors.

The speedup achieved with this strategy, in comparison to a naive parallel iteration over the T^{L2P} ’s targets, increases with the number of calculations in each thread and hence the expansion order p . In an experiment with 32768 leaves, the maximum number of points per leaf, $n_{crit} = 150$, and expansion order $P = 10$, our strategy is 13 % faster. This corresponds to a realistic FMM problem with approximately $1e6$ randomly distributed particles.

Due to their large interaction lists, the previous strategy is too expensive in terms of memory for the near field, T^{M2L} and T^{M2P} operators. For example, allocating an array large enough to store the maximum possible number of source particle

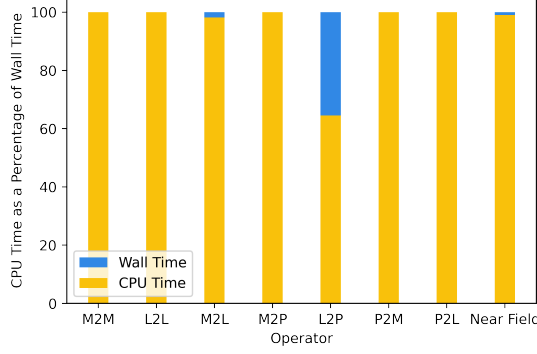


Figure 2.4: CPU time as a percentage of wall time for operators. CPU time is defined as the time in which the algorithm runs pure Numba compiled functions. Wall time is CPU time in addition to the time taken to return control to the Python interpreter, adapted from [31].

coordinates in double precision for the T^{M2P} operator; with $|W| = 148$ and $n_{crit} = 150$, requires $\sim 17\text{GB}$, and a runtime cost for memory allocations that exceeds the computation time. Instead, for the T^{M2L} we perform a parallel loop over the target nodes at each given level, and over the leaf nodes for the M2P and near field, looking up the relevant data from the linear tree as needed. The T^{P2L} interaction list of each target is at most 19 nodes, and the P2M must also calculate a check potential, cancelling out any speedup from cache locality for these operators.

The matrices involved in the T^{M2M} and T^{L2L} operators can be precomputed and scaled at each level [65], and their application is parallelized over all nodes at a given level. Multithreading in this way means that we call the P2P, T^{P2M} , T^{M2P} , T^{L2P} and near field operators once during the algorithm, the T^{M2L} and T^{L2L} are called $d - 2$ times, and the T^{M2M} is called d times, where d is the depth of the octree. This is the minimum number of calls while keeping the operator implementations separate for unit testing.

Figure (2.4) compares the time spent within each Numba-compiled operator (‘CPU time’) to the total runtime (‘wall time’) of each operator. The results are computed over five trials over 32768 leaves, with $n_{crit} = 150$ and $P = 6$, for a random distribution of $1e6$ charges distributed on the surface of a sphere representing a typical FMM problem. The mean size of the interaction lists are $|U| = 11$, $|V| = 42$, $|X| = 3$, $|W| = 3$, and the entire algorithm is computed in $5.95 \pm 0.02s$, with an additional $9.00 \pm 0.01s$ for operator pre-computations for a given dataset, which is unachievable in ordinary single-threaded interpreted Python.

The wall time includes the time to (un)box data, organize inputs for Numba compiled functions, and pass control between Numba and Python. Except for the T^{L2P} which has a different parallelization strategy that requires significant data organization that must take place within the GIL restricted Python interpreter, the runtime costs are less than 5 % of each operator’s total wall time, implying that we are nearly always running multithreaded code and utilizing all available CPU cores.

If Not Python, Then What?

Achieving optimal multithreaded performance with Numba requires careful consideration of the algorithm being accelerated, details of Numba’s backend implementation, as well as a design that suits Numba’s data oriented framework. The pitfalls

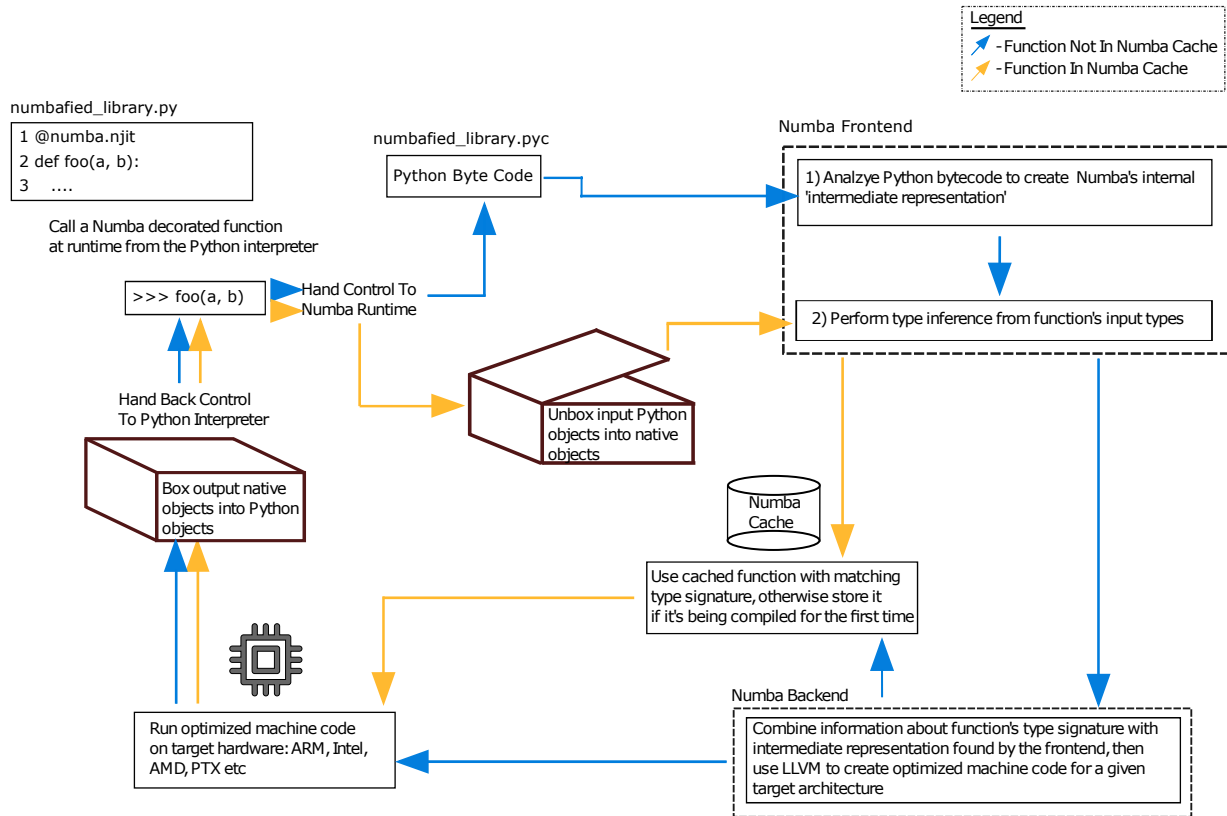


Figure 2.5: Simplified execution path when calling a Numba compiled function from the Python interpreter. The blue path is only taken if the function hasn't been called before. The orange path is taken if a compiled version with the correct type signature already exists in the Numba cache, adapted from [31].

illustrated above show how a user must potentially adapt their code significantly in order to achieve the best performance. Altogether, Numba accelerated code may look decidedly un-Pythonic despite using Python syntax. The complexities involved when using Numba to implement a non-trivial algorithm contrast with its advertisement as a simple way of injecting performance into Python code by applying a decorator. Significant software development expertise is needed in order to optimise a Numba implementation, and arguably more than many in its intended target audience can be expected to possess.

Despite this, Numba is a remarkable tool. Projects which value Python's expressiveness, simple cross platform builds, as well as large open source ecosystem, and only contain a small number of isolated performance bottlenecks would benefit the most from a Numba implementation. Indeed, by writing only in Python our project size is kept minimal with the entire project running to just 4901 lines of code. Furthermore, we are able to deploy PyExaFMM cross platform trivially with Conda and distribute our software in popular Python channels.

Resultantly, we wish to find a middle way, that would retain the usability of a higher level language, with the performance benefits of writing in a lower level language, filling the gap from the handoff between a high-level language interface and the programming constraints, as well as the handoff performance hit of a just in time compiler. We believe that this is offered by Rust, a low-level systems programming language, that retains many of these features.

Listing 2.1: An example of parallel multithreading.

```

import numba
import numpy as np

@numba.njit(cache=True, parallel=True)
def multithreading(A):
    # consider an A as an nxn matrix
    # This is a situation that can lead
    # to thread oversubscription unless
    # Numpy is configured to run single
    # threaded.
    for _ in numba.prange(10):
        B = A @ A

```

Listing 2.2: An example of using Numba in a Python function operating on ndarrays.

```

import numba
import numpy as np

# optionally the decorator can take
# the option nopython=True, which
# disallows Numba from running in
# object mode
@numba.jit
def loop_fusion(A):
    """
    An example of loop fusion, an
    optimization that Numba is able
    to perform on a user's behalf.
    When it recognizes that they are
    operating similarly on a single
    data structure
    """
    for i in range(10):
        A[i] += 1

    for i in range(10):
        A[i] *= 5

    return A

```

Listing 2.3: Three ways of writing a trivial algorithm that passes around a vector, while performing some computations.

```

import numpy as np
import numba
import numba.core
import numba.typed

# Initialize in Python interpreter
data = numba.typed.Dict.empty(
    key_type=numba.core.types.unicode_type,
    value_type=numba.core.types.float64[:]
)

data['v'] = np.ones(100)

# Functions marked with 'njit' rather than
# 'jit' decorator, to force Numba to run in
# no Python mode, disallowing the compilation
# of Python code it is not specialized for.

```

```
# Subroutine 1
@numba.njit
def step_1(data):
    # An example allocation & calculation
    a = np.random.rand(100, 100)
    b = a @ a
    data['step_1'] = data['v']

# Subroutine 2
@numba.njit
def step_2(data):
    # An example allocation & calculation
    A = np.random.rand(100, 100)
    B = A @ A
    data['step_2'] = data['step_1']

@numba.njit
def algorithm1(data):
    # Pays the un(boxing) cost to exchange
    # data with interpreter
    step_1(data); step_2(data)

@numba.njit
def algorithm2():
    # Only pays the boxing cost to return a
    # Python type.
    data = dict()
    data['v'] = np.ones(100)
    step_1(data); step_2(data)
    return data

@numba.njit
def algorithm3():
    # Numba interprets subroutines as a
    # part of a single function body.
    data = dict()
    data['v'] = np.ones(100)

    def step_1(data):
        A = np.random.rand(100, 100)
        B = A @ A
        data['step_1'] = data['v']

    def step_2(data):
        A = np.random.rand(100, 100)
        B = A @ A
        data['step_2'] = data['step_1']

    step_1(data); step_2(data)
    return data
```

2.3 Rust for High Performance Scientific Computing

Interpreted languages, such as Python, Matlab or Julia, offer great usability benefits. They each support large ecosystems of numerical libraries, and offer rapid prototyping due to their interpreted nature. The tools to build software with these languages is designed for portability across a large variety of platforms, from different hardware architectures to operating systems. However, as we have observed in section 2.2, this comes with the overhead of having to deal with an interpreter, and impacts the kind of software that can realistically be built with interpreted languages.

For problems in which this overhead is untenable, compiled languages, such as C, C++ and Fortran are preferred. These languages require greater software engineering expertise, as developers are responsible for allocating memory, installing third party libraries, as well as building their software to target different hardware. These languages in a sense allow a developer to ‘do anything’, of course only if one knows how to do it. The higher software engineering barrier manifests in a dizzying array of compilers, build systems, package managers, testing and documentation libraries and code organisation techniques. The notable thing about all of these optional choices is that it is relatively unclear which is the *preferred* way of doing things, with novice developers likely to find a development strategy that ‘simply works’ and stick to it. Rust stands in contrast to these traditionally preferred compiled languages for high performance scientific computing. Although it is comparably fast it is relatively *inflexible*, with a strongly preferred way of organising, testing, documenting and deploying software. This leads to significantly more uniform Rust code across projects, and a steep, though relatively shorter, learning curve. Importantly, this uniformity makes Rust code easier to share, and port to different operating systems and hardware targets.

Package Managers and Build Systems

For simple programs, it’s tempting to use a compiler directly to create an executable, as in listing (2.4). However, as a project’s codebase expands, with files defined in multiple directories and calls to external libraries often written in other programming languages, this simple one line appeal to a compiler will no longer be sufficient. One could always download their requisite software, and install globally over the machine using a system package manager, and attempt to recompile. However, this quickly becomes untenable if one is developing for a range of hardware and operating system targets, or one needs to use different versions of external binaries and libraries. Therefore, software is usually constructed using a ‘build system’. This is catch all term that refers to a program that takes source files as input and produces a deployable set of binaries or libraries as an output. Classically, build systems were based on ‘Make’ and ‘Autotools’, these softwares generate ‘Makefiles’ - which are recipes for constructing a piece of software given a set of source files and external dependencies. These are robust tools, but the onus is squarely with a developer for ensuring that all dependent libraries are visible to Make, and that the external packages are all of compatible versions.

To manage this complexity, builds are often defined using a metabuild system, most commonly CMake. CMake is a scripting language, and as a meta build system it takes a specification of local and third party dependencies and hardware targets,

and generates Makefiles. CMake gives developers a great deal of flexibility, it is multi-platform, and language agnostic, however using it directly is not straightforward. Indeed, there is a significant body of literature discussing best practices with CMake [54], however CMake is not responsible for downloading and installing third party packages, or verifying their relative compatibility, implementing its best practices are again left to users.

Approaches for reliable builds vary amongst projects, ranging from ‘low tech’ readme driven solutions, in which a user follows a recipe of instructions, to more modern automated solutions. Figure (2.6) provides an overview of the different approaches taken. The lack of a uniform standard means that some projects go as far as to implement a custom build system, such as the Boost library [4].

Keeping on top of packages, which are constantly being iterated upon independently, and may support different features with each release, is nearly impossible to do manually. Furthermore, one may want to support a specific set of packages, and respectively versions, for a given hardware or operating system target, but a different set for a different build of the same software. This has led to the development of ‘package managers’ which are a catch all term for softwares that can download and install required software for you, and often verify whether version constraints are satisfied too. Package managers can be delineated into ‘system package managers’, which download operating system specific packages written in any language globally on your machine, and ‘language specific package managers’ which focus only on packages written in a given target language, but are operating system agnostic from a user’s perspective.

In terms of system package managers Linux examples include apt, yum and debian, for MacOS there is Homebrew, and Chocolatey for Windows. These can handle both source installations, in which source code is compiled upon download, as well as binary installation, in which pre-built binaries matching your hardware constraints are installed. Binary installation is often preferred, as it is faster, and often reflects a stable release. Language specific package managers, such as pip for Python, which can handle dependencies written in these languages only. Developing your own packages for system package managers is not developer friendly, official package repositories of Linux package managers are moderated by their respective maintainers, though it is possible to set up a personal repository this is quite a sophisticated approach for simple research outputs. The simple alternative, pursued by up to 26 % of surveyed C++ softwares (see fig. (2.6)), is to simply add third party software as a submodule. Often to ease installation, C++ libraries are shipped as ‘header only’ libraries, this makes them easy to install, requiring only an `#include`. However by using such techniques a project can quickly grow to contain thousands, to millions of lines of code, much of it in dependencies. Compiling from source has the potential to be a painstakingly slow process, and must be repeated for every combination of hardware and operating system targets one wants to run on.

With the exception of Fortran, which has made recent strides to develop a standardised modern package manager and build system, inspired by Rust’s Cargo [19], C and C++ do not have a single officially supported package manager or build system. The resulting landscape is a multitude of package managers [55, 63, 12] and build systems [61, 5, 53] a few of which we have cited here, all of which replicate each others functionality, none of which are universally accepted or implemented across projects nor officially supported by the C++ software foundation.

Listing 2.4: Compiling a simple `source_code.cpp` into a `compiled_binary` file from a terminal.

Installing and Building Software in C++/Fortran

Builds for open source software are often **Readme Driven**. In this common approach developers provide a set of instructions for how to install a project's dependencies and the project itself. Common approaches include:

Dependency Management Methods



1) Simply add all dependencies to source tree of your project. Projects with a large number of dependencies can grow to have millions of lines of code, which can have a drastic effect on compilation times. Large C++ projects for example can take between several minutes to several hours to compile from source



2) Use a system package manager, and source from repositories such as GitHub to install globally. These can then be found by build systems. This makes it difficult to build isolated build environments, and configure builds with different versions, or sets, of dependencies.

Build Methods



1) Developer provided Make and Autotools scripts. lowest barrier to entry but build system will use globally installed dependency libraries, unless alternative provided. Makes multi-platform builds, or those using different software versions challenging.

2) Developer provided CMake scripts, to generate build systems for different environments. Again, reliant on globally installed dependency libraries.



Neither of these methods can check for dependency conflicts, which is left up to the developer to resolve.

Modern Package Management

Modern package managers allow for maximum safety and flexibility. They support multiple operating systems, compilers, build systems, and hardware microarchitectures, and are usually defined by recipes written in a simple scripting language such as Python, and can be used to generate build system scripts such as Makefiles and CMake scripts. They also support dependency tree checking for conflicting requirements, leading to stable builds. Two popular leading tools for HPC in C++ and Fortran are **Spack** and **Conan**.



Spack

- + Supported on Linux and MacOS.
- + Package recipes specified with Python scripts.
- + Over 5000 commonly used packages available.
- + Supports isolated environments, supporting building packages with from specific constraints, and clean global environment.
- + Can target hardware microarchitectures.
- + Growing support for binary package installation, though not available on all hardware targets.
- + Compatible with all major build systems, such as CMake and SCons
- No Windows support.



Conan

- + Supported on Linux, MacOS and Windows
- + Package recipes specified with Python scripts.
- + Supports isolated environments, supporting building packages with from specific constraints, and clean global environment.
- + Support for binary packages, speeding up installation times.
- + Compatible with all major build systems, such as CMake and SCons.

However, we note that even in 2021 modern package managers are still only used in a tiny fraction of all projects. For example Conan, only accounts for 5% of all C++ projects surveyed by JetBrains¹ in comparison to 21% of projects still relying on the system package manager, 26% simply including a dependencies source code as a part of a project's source tree, 23% using the readme driven build of each specific dependency, and 21% installing non-optimised precompiled binaries from the internet. Indeed only a small minority, 22%, used a package manager of any kind, with no solution taking a majority of even this market share. This is in stark contrast to the uniformity of the situation in Rust, in which all projects use Cargo as a build system and package manager and rustc as a compiler.

1. <https://www.jetbrains.com/lp/devecosystem-2021/cpp/>

Figure 2.6: An overview of building software in other compiled languages.

```
>>> gcc -o compiled_binary source_code.cpp
>>> # Can then run with
>>> ./compiled_binary
```

Recent years have seen the bundling of package managers *with* build systems, resulting in softwares that can simply take source files, and a set of dependencies, and resolve a single binary output for a given hardware and software target. Examples include Conda for Python, and Rust’s Cargo system. These modern efforts are able to compile both source and binary packages for all operating systems, and can target a wide range of hardware architectures. Furthermore, the trend has been towards the specification of project dependencies in a single structured text file (see listing (2.5)), which is then handled by the package manager with no further user effort. Efforts have been in this direction for older compiled languages, the most notable examples being Spack and Conan (see fig. (2.6)). Importantly, in the case of Rust, Cargo is officially supported and shipped as a core part of its runtime. In contrast to C, C++ and Fortran, Rust has a *single* officially supported compiler, `rustc`. By taking global decisions for all software written in Rust, a significant burden is removed from developers of Rust software, similar to the situation in many interpreted languages. Indeed Cargo builds are often executed in a *single line*, eradicating the complex readme driven builds common in other compiled languages. For the project specified by listing (2.5) can be compiled from a terminal with the command: `cargo build`.

Listing 2.5: An example of a simple `Cargo.toml` file for a Rust project, with source as well as binary dependencies.

```
[package]
name = "new-package"
version = "1.1.0"
authors = ["Foo Bar"]
edition = "2018"
description = "A great new package"
license = "BSD-3-Clause"
homepage = "https://github.com/foo_bar/new-package"
repository = "https://github.com/foo_bar/new-package"
keywords = ["numerics"]
categories = ["mathematics", "science"]

[dependencies]
# Build from binary on crates.io
memoffset = "0.6"
rand = "0.8.4"
itertools = "0.10"
vtkio = "0.6"

# Build from source on GitHub
mpi = { git = "https://github.com/rsmapi/rsmapi", branch = "main" }
```

Code Organisation and Quality

Rust introduces many ergonomic features for code organization. The most novel features, which may not be familiar to those coming from other compiled languages, are the concepts of traits, lifetimes and the borrow checker.

Traits

Traits are Rust’s system for specifying shared behaviour and building abstraction, constituting a wholesale replacement of object oriented programming, with its inheritance based hierarchies. Traits only enforce *behaviour*, and therefore are strictly less brittle than object orientation which enforces a type. We provide some example syntax in listing (2.6), and contrast it to equivalent object oriented code in ;listing (2.7). We notice that the object oriented code has a built in hierarchy, which means that adding shared behaviour will effect our **MyType** type, and everything that subsequently depends on it, in contrast to the trait based code in which we can inject new behaviour into MyType, its subsequent dependents we only need to know about the traits and their associated interfaces that they themselves rely on. This means that one can focus solely on the behaviour implicit in a given Trait, rather than having to comb through a potentially complex hierarchy of objects and inheritance to understand what a given line of code is doing, making large Rust projects significantly more readable than their object oriented equivalents.

Traits can be seen to specify shared behaviour in a bottom up manner, as opposed to top down object orientation, new features can be injected into existing types. This is useful for scientific programming, as we are usually concerned with data oriented programming. As we saw in section 2.2 exploring data oriented programming in Python, object oriented design obfuscates operations on the data itself behind abstraction.

Listing 2.6: An example of a trait implemented on a custom type.

```
// Custom type declaration , with its own set of attributes
// and methods
pub struct MyType {
    pub attribute String;

    pub fn bar() {
        println!("I'm a default implementation defined on this type.")
    }
};

// Traits specify an interface
pub trait MyTrait {
    fn foo(&self) -> String;
};

pub trait MyOtherTrait {
    fn baz(&self) -> String;
}

// We can implement this interface for our type
impl MyTrait for MyType {
    pub fn foo() {
        format!("I'm a required method for this trait.")
    }
};

impl MyOtherTrait for MyType {
    pub fn baz() {
        format!("I'm a required method for this trait.")
    }
}

fn main () {
```

```
let mytype = MyType{attribute: "boz"};

// Use behaviour from MyTrait
println!(mytype.foo());

// Use behaviour from MyOtherTrait
println!(mytype.baz());

// Use behaviour defined by MyType itself
println!(mytype.bar());
}
```

Listing 2.7: An example of behaviour in listing (2.7) implemented in an object oriented manner

```
from abc import ABC, abstractmethod

class MyTrait(ABC):
    def __init__(self):
        pass

    @abstractmethod
    def foo(self):
        """A required method for this class"""
        pass

class MyOtherTrait(ABC):
    def __init__(self):
        pass

    @abstractmethod
    def baz(self):
        """A required method for this class"""
        pass

class MyType(MyTrait, MyOtherTrait):
    def __init__(self, attribute):
        self.attribute = attribute

    def bar(self):
        print("I'm a default implementation defined on this type.")

    def bar(self):
        print("I'm a required method for the MyTrait class.")

    def baz(self):
        print("I'm a required method for the MyTrait class.")

def main():
    mytype = MyType("boz")

    # Use behaviour from MyTrait
    print(mytype.foo())

    # Use behaviour from MyOtherTrait
    print(mytype.baz())

    # Use behaviour defined by MyType itself
    print(mytype.bar())
```

This isn't to say similar syntax isn't available in other compiled languages. In C++20, 'concepts' were introduced as a trait like mechanism to specify interfaces. However, as with many things in C++, this isn't enforced and it is up to a user to choose to implement their code in this manner. This is borne out in the fact that concepts are *nominally typed*, and therefore don't enforce behaviour as in Rust's *structurally typed* traits. The implication of this being that a type may 'accidentally' implement a concept, if it happens to define its relevant methods. A consequence of this is that a valid C++ program can contain a confusing mix of trait based, and object oriented code, with a reader then dependent on documentation to understand a software's behaviour, and what has been enforced and where, by the developer.

Lifetimes and the Borrow Checker

Another new Rust concept for developers coming from other compiled languages is the idea of 'lifetimes' and 'ownership'. Every reference in Rust has an associated lifetime, and a singular 'owner'. Which enforce the programming pattern of 'resource allocation is initialisation' (RAII) as a feature of the Rust compiler. The basic rule is that references are owned within a scope, and dropped when out of scope. We provide a basic example in listing (2.8), in a Rust context RAII is better encapsulated by another acronym 'destructors run at exit scope' (DRES).

Listing 2.8: A demonstration of basic rules of borrowing and ownership.

```
fn create_vec() -> Vec<i32> {
    // Create a vector, and pass ownership to caller
    let mut vec = Vec::new();
    let n = 5;
    for i in 1..n {
        vec.push(i)
    }
    vec
}

fn processor_borrows(vec: &Vec<i32>) {
    // Function doesn't take ownership of vector

    for i in vec.iter() {
        println!("i: {}", i);
    }
}

fn processor_borrows_mut(vec: &mut Vec<i32>) {
    // Function doesn't take ownership of vector,
    // but can mutate it via a mutable reference

    for i in 0..vec.len() {
        vec[i] = vec[i]*2;
    }
}

fn processor_owns(vec: Vec<i32>) {

    for i in vec.iter() {
        println!("i: {}", i);
    }
}
```

```
fn main() {
    let vec = create_vec();
    processor_owns(vec);

    // Calling functions below now will cause an error,
    // as ownership of vec has been passed
    processor_borrows(&vec);
    processor_borrows_mut(&mut vec);
    // to 'processor_owns', resulting in error message:
    // 43 |         processor_borrows(&vec)
    //      ^^^^^ value borrowed here after move

    // The below is valid, as there is still only one mutable reference
    let mut vec2 = create_vec();
    processor_borrows_mut(&mut vec2);
    processor_borrows(&vec2);

    // The following will not work, as we create two mutable references
    let r1 = &mut vec2;
    let r2 = &mut vec2;
    processor_borrows_mut(r1);
    processor_borrows_mut(r2);
    // obtaining the following error:
    // 53 |         let r1 = &mut vec2;
    //      |         _____ first mutable borrow occurs here
    // 54 |         let r2 = &mut vec2;
    //      |         ^^^^^^^^^ second mutable borrow occurs here
}
```

Lifetimes are a Rust concept that guarantees that any references to a resource live at least as long as the resource itself. The main aim of this to prevent dangling references, whereby references to unintended data, or de-allocated data, are not present in the compiled binary. This removes a large class of common bugs from Rust software such as dangling pointers, and double free errors. We illustrate this in listing (2.9), which won't compile as the inner scope defines a value **x**, which doesn't survive into the outer scope.

Listing 2.9: A demonstration of lifetimes as a function of scope. This example is adapted from the Cargo book [32].

```
fn main() {
    let r;

    // Lifetimes are defined by scope.
    {
        let x = 5;
        r = &x
    }
    println!("r: {}", r);
    // Results in error
    // 6 |         r = &x;
    //    |         ^^ borrowed value does not live long enough
    // 7 |     }
}
```

Rust's compiler enforces ownership and lifetime rules uses a program called the 'borrow checker' to ensure that all references, and lifetimes, are referring to valid memory locations at *compile time* without the need for a runtime garbage collector.

This is a huge advantage over other compiled languages, for example memory related bugs have been found to constitute as much as 70 % of all security bugs at Microsoft [43]. From a scientific programming perspective, handling memory additionally bogs down algorithm development, iteration, and ultimately publication, and was one of the major contributing factors in the push towards interpreted languages in recent decades. Again, this is not to say similar features do not exist in other compiled languages. C++ has its notion of ‘smart pointers’, which enforce RAII principles too, however as with other features in C++, this is a library feature rather than a core part of its compiler, or language specification, making it optional and unenforced.

The benefits of Rust’s memory system extend to parallel code too. A wide variety of programming paradigms for parallel computation with threads exist, in scientific software we often use OpenMP, which uses the shared memory paradigm, such that all threads operate on the same data. In this setting one of the most common bugs is a data race condition, whereby multiple threads attempt to write to the same memory location. Rust’s ownership system enforces the atomicity of all memory operations, by passing ownership between threads, as only a single mutable reference can exist to a piece of memory Rust is able to prevent data races at compile time. Although atomic operations are available in OpenMP, or other shared memory systems, the benefit of Rust is that we know that our compiled code *cannot* contain this bug, without having to rely on unit tests for expected behaviour.

Code Quality

Rust’s runtime includes a test runner, a documentation generator, and a code formatter. As with other Rust features, these are maintained in lock step with the language specification, and with reference to other Rust developments. This imposes universal constraints on all Rust projects, allowing for objectively defined ‘good’ Rust code, rather than relying on various standards of best practices that vary between projects and organisations.

Rust’s Scientific Ecosystem and Foreign Libraries

Despite being a young language, Rust already supports a mature ecosystem of libraries for scientific computing with high-level multithreading support [57], numerical data containers [46], and tools for generating interfaces to Python [41].

Many tools are yet to be ported into native Rust, however high quality bindings exist for core tools such as MPI [56], BLAS and LAPACK [6]. The problem with interfacing with tools written in other languages is again related to building software, however Cargo offers tools to build software written in other languages and integrate it with Rust code via the **build** crate, which allows one to leverage existing build systems written for software written in foreign languages. This detracts from the benefits offered by Cargo as a unified package manager and build system, raising similar problems to those encountered when building software in other compiled languages. However, we observe that this remains a concern of the software’s developer, who is responsible for providing build scripts for the operating systems and hardware platforms that they wish to support, and from a downstream user perspective their build process remains the same as with pure Rust packages, where the dependency is defined in their **Cargo.toml** file.

We also note that Rust is missing key tools for scientific computing, such as a code generation for GPUs, as well as for advanced optimised advanced linear algebra

routines, especially for sparse matrices. However, both of these applications are an active area of development.

Conclusion

Rust’s syntax and program structure are strongly opinionated, its build system emphasises simplicity and portability, and its ecosystem for scientific computing is rapidly developing. A small software team, as is common in academic settings, writing in Rust are effectively able to maintain and deploy Rust projects to a wide variety of platforms, from laptops to supercomputers. Rust’s simplified build process results in minimal configuration for users, encouraging the widespread adoption of Rust projects.

2.4 Case Study: RustyTree, a Rust based Parallel Octree

Octrees are a foundational data structure for fast algorithms in \mathbb{R}^3 . Writing software for octrees that can be distributed across parallel computing systems is challenging, demonstrated by the limited availability of open-source software [9, 15]. This is inspite of their ubiquity in scientific computing applications from adaptive finite-element methods to many-body algorithms [60]. We present Rusty Tree, an implementation of MPI-distributed octrees with a complete Python interface. Rusty Tree is a proof of concept for Rust as a tool for high-performance computational science. We document our experience in developing Rusty Tree, and use it as a tool to explore the current landscape of scientific computing with Rust. From multithreading, and developing a Python interface, to distributing applications with MPI, writing Rusty Tree involved using much of Rust’s scientific computing ecosystem.

Parallel Octrees

As mentioned in section 2.2, there are multiple representations of octrees. We choose a linear octree representation, in which we discard interior nodes and store a list of leaves, as in our previous work PyExaFMM [31]. These lie in contrast to pointer based octrees in which each node stores pointers to its parent as well as children, which have a relatively higher storage cost, in addition to a synchronisation and communication overhead for parallel implementations that must keep track of pointers across nodes [62].

Pointer based octrees are usually constructed ‘top down’. A user specifies a threshold for the maximum number of particles in a leaf node, n_{crit} , and a unit cube encapsulating the problem domain is refined until this is satisfied. After refinement, one can optionally ‘balance’ adjacent leaf nodes such that adjacent nodes obey a relative size constraint. However, translating this logic to a parallel setting is challenging. Constructing octrees ‘top down’ involve significant data communication. As the point distribution on each processor isn’t known a priori, each processor is responsible for constructing its own octree from a random batch of points, reconciling the trees across nodes then requiring a global parallel merge. Furthermore, there are likely to be significant overlaps and duplicates in the nodes of each processor’s octree, which would result further refinement after the parallel merge in order to




Figure 2.7: How to Morton encode a node in an octree. The node is described by an ‘anchor’ coordinate.

satisfy the n_{crit} constraint. Balancing the octree, as well as load-balancing the work assigned to each processor, must be done as post-processing steps.

For these reasons, Sundar et al introduced ‘bottom up’ tree construction [60]. ‘Bottom up’ refers to the fact that octrees are constructed from the leaf level upwards, which allows us to incorporate the initial data distribution into our algorithms. In order to do this they use a space-filling curve that passes through each octant, for which they choose a Morton encoding.

Given a set of distributed point data spread across n_p processors, our strategy is as follows:

1. Generate a Morton key for each point at the leaf level corresponding to a user defined octree depth.
2. Perform a parallel sort of these leaf octants using MPI, such that processors with adjacent MPI ranks have adjacent data.
3. *Linearise* leaves on each processor, such that their leaves do not overlap, in a manner that favours smaller leaves over larger ones, as in algorithm (4).
4. *Complete* the space defined by the leaves on each processor, and find the domain occupied by their leaves, see figure (2.8) and algorithm (5). Define the coarsest nodes of this set as *seeds*.
5. Construct a *block tree*, by completing the space between the seeds with the minimum spanning set of nodes, described in algorithm (7). The nodes of this block tree are referred to as *blocks*. This gives us a distributed coarse complete linear octree that is based on the underlying data distribution.
6. Point data is communicated to the processor containing its associated block, and the blocks are refined based on n_{crit} , providing the final linear octree tree.
7. Optionally, the tree is *balanced*. This can be done for the local subtree on each processor using algorithm (6). The locally balanced trees can then be sorted in parallel again, and locally linearized as in step (3), to provide a globally balanced tree.

This approach is summarised in algorithm (2). The complexity of this process is bounded by the parallel sorts, which for randomly distributed point data, run in $O(N_{\text{leaf}} \log(N_{\text{leaf}}))$ work and $O(\frac{N_{\text{leaf}}}{n_p} \log(\frac{N_{\text{leaf}}}{n_p}) + n_p \log(n_p))$ time where N_{leaf} is the number of leaves in the final octree, and n_p is the number of processors (see appendix A.2 for a derivation). This makes the efficiency of the parallel sort a bottleneck in tree construction.

Algorithm 2 Construct Distributed Octree (Parallel)

Input: A distributed list of points L , and a parameter n_{crit} specifying the maximum number of points per octant.

Output: A complete linear octree, B .

Work: $O(n \log n)$, where $n = \text{len}(L)$.

Storage: $O(n)$, where $n = \text{len}(L)$.

$F \leftarrow [\text{Octant}(p, \text{MaxDepth}), \forall p \in L]$

$\text{ParallelSort}(F)$

$B \leftarrow \text{BlockPartition}(F)$, using algorithm (8)

for $b \in B$ **do**

if $\text{NumberOfPoints}(b) > n_{crit}$ **then**

$B \leftarrow B - b + \text{Children}(b)$

end if

end for

Optional Balancing over subtrees, f .

if $\text{Balance} = \text{True}$ **then**

for $f \in F$ **do**

$\text{Balance}(f)$, using algorithm (6)

end for

$\text{ParallelSort}(F)$

for $f \in F$ **do**

$\text{Linearise}(f)$, using algorithm (4)

end for

end if




Figure 2.8: Adapted from [60], describing the algorithms in \mathbb{R}^2 for clarity. Figure (b) shows the minimal tree constructed using algorithm (5) where the two octants in (a) are used as an input. Figure (c) shows a selection of *seeds*, which are used as an input to algorithm (7), resulting in (d). Figure (d) also shows two octants generated to complete the domain. The first one is the coarsest ancestor of the least input octant which does not overlap with the least octant in the input. The second is the coarsest ancestor of the greatest possible octant, which doesn't overlap with it.

We extend the approach first presented by Sundar et al in two key ways. Firstly, we use HykSort [59], a variant of parallel hypercube sort, which reduces global data communication cost in comparison to a parallel sample sort based on **MPIAllToAllv**. We cover the relative merits of HykSort in the next subsection. Secondly, balancing algorithms in the original implementation are based on a ‘ripple propagation scheme’. The ‘ripple effect’ is the phenomenon that describes how octants can effect immediately adjacent octants if they are split, or combined, which can trigger further changes, that must then be propagated across the network. These algorithms involve relatively complex communication schemes in their implementation. We therefore replace this with an algorithm which, starting with an unbalanced tree, finds the locally balanced tree on each processor using algorithm (6), performs another parallel sort using HykSort, and removes overlaps, favouring smaller octants, using algorithm (4). This guarantees that the balance condition is maintained. This logic will recover a globally balanced tree, with a significantly simpler communication pattern. This strategy has been previously been implemented in large scale parallel octree construction [39], and has been shown to obtain similar runtimes to calculate balancing when compared to ripple propagation based schemes [58].

Load balancing across processors, while maintaining Morton order is challenging. Currently, the data partitioning imposed by the parallel sort imposes a degree of load balancing, which we examine for different geometries in figure (2.12c). However, in future we would like to extend our work to incorporate an explicit load balancing step. Our current idea is based on a sample sort, with splitter selection and bucket sizes chosen to reflect the mean load we wish to assign to each processor.

Algorithm 3 Morton Order (Sequential): Find the lesser of two Morton keys with different anchors.

Input: Two Morton ids, A and B , with different anchors.

Output: R , the lesser Morton key.

$X_i \leftarrow (A_i \oplus B_i), i \in x, y, z$

$e \leftarrow \arg \max (|\log_2(X_i)|)$

if $A_e < B_e$ **then**

$R \leftarrow A$

else

$R \leftarrow B$

end if

Algorithm 4 Remove Overlaps From Sorted List of Octants (Sequential)

- **Linearise.** Favour smaller octants over larger overlapping octants.

Input: A sorted list of octants, W .

Output: R , an octree with no overlaps.

Work: $O(n)$, where $n = \text{len}(W)$.

Storage: $O(n)$, where $n = \text{len}(W)$.

for $i \leftarrow 1$ to $\text{len}(W)$ **do**

if $W[i] \notin \{\text{Ancestors}(W[i+1]), W[i+1]\}$ **then**

$R \leftarrow R + W[i]$

end if

end for

$R \leftarrow R + W[\text{len}(i)]$

Optimised Morton Encoding

Lookup Tables

Consider a given set of indices describing the anchor of an octree node (x, y, z) . The Morton encoding can be statically encoded in a lookup table.

For example, $(x, y, z) = (4, 55, 132)$ which in binary is $(100, 110111, 10000100)$. As we have to interleave these bits, we can add zero bits and perform a bitwise 'or' operation to find the final Morton key:

```
x_shift | y_shift | z_shift =
01000000
| 00010010000010010010
| 100000000000000100000000
```

If $x, y, z \in [0, 255]$, these shifts can be stored in small statically stored lookup tables, and we can encode a Morton key for larger indices by considering their bits byte by byte.

This is a divide and conquer strategy, that has been shown to be faster than on-the-fly implementation using for-loops and bit-shifts¹.

1. <https://www.forceflow.be/2013/10/07/morton-encoding-decoding-through-bit-interleaving-implementations/>

Shared Memory Parallelism With Rayon

This lookup strategy can be performed in parallel over each anchor being encoded, which is easy to do using the Rayon crate for shared memory parallelism.

Rust defines 'iterators', which are a functional programming abstraction to apply a single transformation to a set of data.

For example, consider the calculation of the squares of a vector of numbers, which are then summed. This can be expressed as:

```
let vec: Vec<i32> = vec![0,1,2,3,4,5];

let res = vec.iter()
    .map(|&i| i*i)
    .sum();
```

Rayon's main abstraction is to extend this to a parallel setting with *parallel iterators*:

```
let res = vec.par_iter() // <- only change
    .map(|&i| i*i)
    .sum();
```

Rayon's parallel iterators are an abstraction built on top of its work-stealing based parallel backend. Rayon fully incorporates Rust's multithreading safety features, and therefore guarantees data-race freedom.

Figure 2.9: Optimising Morton encodings using shared memory parallelism and lookup tables.

Algorithm 5 Construct a Minimal Linear Octree Between Two Octants (Sequential) - CompleteRegion.

Input: Two octants a and b , where $a > b$ in Morton order.

Output: R , minimal linear octree between a and b .

Work: $O(n \log n)$, where $n = \text{len}(R)$.

Storage: $O(n)$, where $n = \text{len}(R)$.

```
for  $w \in W$  do
    if  $a < w < b$  and  $w \notin \{\text{Ancestors}(b)\}$  then
         $R \leftarrow R + w$ 
    else if  $w \notin \{\text{Ancestors}(a), \text{Ancestors}(b)\}$  then
         $W \leftarrow W - w + \text{Children}(w)$ 
    end if
end for
Sort( $R$ )
```

Algorithm 6 Balance a Local Octree (Sequential) - Balance. A 2:1 balancing is enforced, such that adjacent octants are at most twice as large as each other.

Input: A local octree W , on a given node.
Output: R , a 2:1 balanced octree.
Work: $O(n \log n)$, where $n = \text{len}(R)$.
Storage: $O(n)$, where $n = \text{len}(W)$.
 $R = \text{Linearize}(W)$
for $l \leftarrow \text{Depth}$ **to** 1 **do**
 $Q \leftarrow \{x \in W \mid \text{Level}(x) = l\}$
 for $q \in Q$ **do**
 for $n \in \{\text{Neighbours}(q), q\}$ **do**
 if $n \notin R$ and $\text{Parent}(n) \notin R$ **then**
 $R \leftarrow R + \text{Parent}(n)$
 $R \leftarrow R + \text{Siblings}(\text{Parent}(n))$
 end if
 end for
 end for
end for

Algorithm 7 Construct a Complete Linear Octree From a Set of Seed Octants Spread Across Processors (Parallel) - CompleteOctree

Input: A distributed sorted list of seeds L .
Output: R , a complete linear octree.
Work: $O(n \log n)$, where $n = \text{len}(R)$.
Storage: $O(n)$, where $n = \text{len}(R)$.
 $L \leftarrow \text{Linearise}(L)$, using algorithm (4).
if rank = 0 **then**
 $L.\text{push_front}(\text{FirstChild}(\text{FinestAncestors}(\text{DeepestFirstDescendent}(\text{root}), L[1])))$
end if
if rank = $n_p - 1$ **then**
 $L.\text{push_back}(\text{LastChild}(\text{FinestAncestors}(\text{DeepestLastDescendent}(\text{root}), L[\text{len}(L)])))$
end if
if rank \neq 0 **then**
 Send($L[1]$, (rank-1))
end if
if rank \neq ($n_p - 1$) **then**
 $L.\text{push_back}(\text{Receive}())$
end if
for $i \leftarrow 1$ **to** ($\text{len}(L) - 1$) **do**
 $A \leftarrow \text{CompleteRegion}(L[i], L[i + 1])$, using algorithm (5)
end for
if rank = $n_p - 1$ **then**
 $R \leftarrow R + L[L]$
end if

Algorithm 8 Partitioning Octants Into Coarse Parallel Blocks (Parallel)
 - BlockPartition.

Input: A distributed list of octants F .

Output: A list of blocks G , F redistributed but the relative order of the octants is preserved.

Work: $O(n)$, where $n = \text{len}(F)$.

Storage: $O(n)$, where $n = \text{len}(F)$.

$T \leftarrow \text{CompleteRegion}(F[1], F[\text{len}(F)])$, using algorithm (5)

$C \leftarrow \{x \in T \mid \forall y \in T, \text{Level}(x) \leq \text{Level}(y)\}$

$G \leftarrow \text{CompleteOctree}(C)$, using algorithm (7)

for $g \in G$ **do** $\text{weight}(g) \leftarrow \text{len}(F_{\text{global}} \cap \{g, \{\text{Descendants}(g)\}\})$

end for

$F \leftarrow F_{\text{global}} \cap \{g, \{\text{Descendants}(g)\} \mid g \in G\}$

Efficient Parallel Sorts

Designing and implementing an efficient sorting algorithm that can scale to thousands of cores is difficult since it requires irregular data access, communication, and equal load-balance. As first presented by Sundar et al, parallel sorts in the creation of octrees are implemented using a variant of Sample Sort [60]. Briefly, this approach samples elements at each processor to create a set of $n_p - 1$ ordered ‘splitters’, which are shared across all processors and define a set of n_p buckets. This is followed by a global all-to-all communication call over all n_p processors to assign elements to their corresponding bucket. Finally, a local sort is performed at each bucket to produce a globally sorted array. SampleSort is well understood. However, its performance is quite sensitive to the selection of splitters, which can result in load imbalance. Most importantly, the all-to-all key redistribution scales linearly with the number of tasks and can congest the network. As a result sample sort may scale sub-optimally, especially when the communication volume approaches the available hardware limits [59].

An alternative approach is provided by HykSort [59], which is a generalisation of Quicksort over a hypercube [64] from 2-way splits, to k -way splits, with the addition of an optimised algorithm to select splitters. Quicksort, Hyksort and Sample Sort are compared in figure (2.10). Instead of splitting the global array into n_p buckets, Hyksort splits it into $k < n_p$, and recursively sorts for each bucket. We notice that at each recursion step, each task communicates with just k other tasks. Indeed, for $k = 2$ we recover Quicksort over a hypercube. Both Hyksort, and the parallel splitter selection algorithm are provided in appendix (A.3), alongside complexity comparisons between Hyksort and Sample Sort. We note that Hyksort has a lower asymptotic complexity, with no terms that scale linearly in number of tasks in the MPI communicator, p , unlike Sample Sort.

Rusty Tree

Developed using the algorithms described in the previous sections, Rusty Tree can be found at <https://github.com/rusty-fast-solvers/rusty-tree>. It is consumed either as a Rust crate or a Python package, with binaries distributed via Conda. We have so far tested Rusty Tree on MacOS as well as various Linux platforms including RHEL, common to many HPC platforms, and Ubuntu. With a fully developed suite of tests, documentation and Python interfaces, Rusty Tree represents

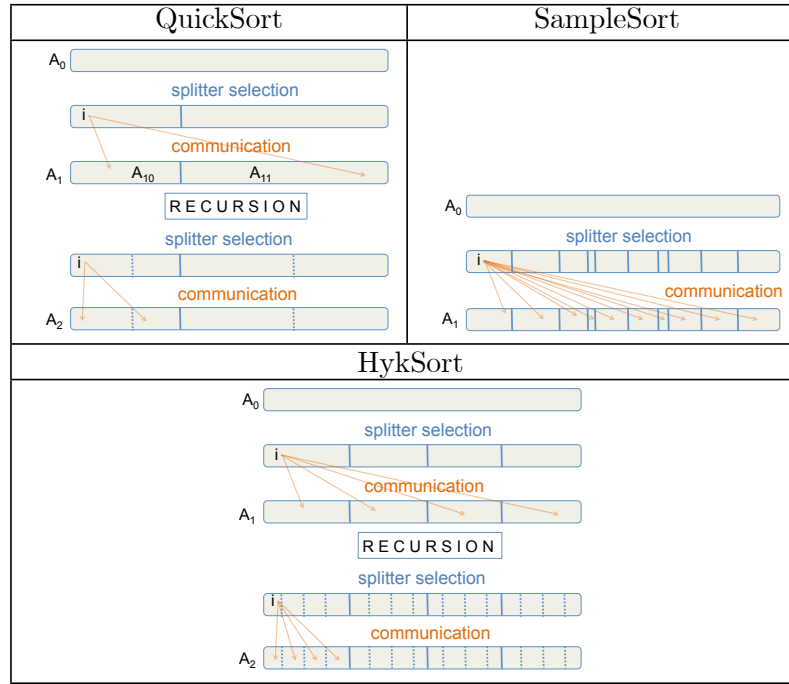


Figure 2.10: Communication pattern of Hyksort algorithm compared to a parallel Sample Sort, as well as a Quicksort over a hypercube, adapted from [59]. We see that HykSort results in a lower communication overhead than Sample Sort.

a model example of our final software infrastructure for fast solvers. Example usage in Python mirrors the Rust interface as shown in listings (2.10) and (2.11). Users are able to use point data generated at runtime, or provide a HDF5 file for Rusty Tree to consume. Rusty Tree is also capable of outputting a VTK file for visualisation. We follow data oriented principles in Rusty Tree’s design, with a shallow abstraction layer that acts as an interface to operations on point data, and encoded Morton indices. An abridged software architecture in UML style is shown in figure (2.11). In terms of lines of code, this translates into just 2578 lines of Rust and 688 lines of Python, including test code, with less than 100 lines of configuration code in terms of Toml and Yaml files specifying the Rust and Python builds.

The Python interfaces are built using Maturin, a tool for building and publishing Rust-based Python packages. It can do this in different ways, including the popular PyO3 [47] and rust-cpython [52] crates both of which offer Rust bindings for Python, allowing one to use Rust libraries as native extension modules as well as interacting with Python code from Rust binaries. However Maturin also supports creating bindings using Rust’s C Foreign Function Interface (CFFI), using the `clib` crate to expose Rust types to consumers of the C application binary interface (ABI). We choose the CFFI, as the types we wish to expose into Rust are relatively simple pointers to Rust iterators and the pointer to an MPI communicator created from Python. This also allows for a simple design with minimal external libraries, with our Python interface to be cleanly separated from Rust source code.

Listing 2.10: Using Rusty Tree from its Python package.

```
from mpi4py import MPI
import numpy as np

from rusty_tree.distributed import DistributedTree

# Setup communicator
```

```

comm = MPI.COMM_WORLD

# Cartesian points at this process
points = np.random.rand(100000, 3)

# Generate a balanced tree
balanced = DistributedTree.from_global_points(points, True, comm)

# Generate an unbalanced tree
unbalanced = DistributedTree.from_global_points(points, False, comm)

# Trees map between points and the keys that they correspond to contain
point = balanced.points[0].head # grab a point
print(balanced[point]) # find the key that it maps

# Trees map between keys and the points they contain
key = balanced.keys[0].head # grab a point
print(balanced[key]) # find the points that this key contains

# Trees implement iterator protocol as well as slicing and indexing
# (without copy of underlying Rust data)
# Slice of 10 keys
key_slice = balanced.keys[:10]

# Slice of 10 points
point_slice = balanced.points[:10]

def foo(key):
    """Some key processor function"""
    pass

# Iterating in this way avoids copying data from Rust to Python
for key in key_slice:
    foo(key)

# Copy only performed when printing in Python
print(point_slice)
print(key_slice)

```

Listing 2.11: Using Rusty Tree as a Rust library.

```

use rand::prelude::*;
use rand::SeedableRng;

use mpi::{
    environment::Universe,
    topology::UserCommunicator,
    traits::*,
};

use rusty_tree::{
    distributed::DistributedTree,
};

fn main () {
    // Generate a set of randomly distributed points
    let mut rng = StdRng::seed_from_u64(0);
    let between = rand::distributions::Uniform::from(0.0..1.0);
    let mut points: Vec<[PointType; 3]> = Vec::new();
    let npoints = 1000000;

```




Figure 2.11: We see that the shallow hierarchy inherent in Rusty Tree’s design, with a ‘Distributed Tree’ interface in both Python and Rust which is a wrapper over functionality that acts directly on data. On the Python side we have an additional **Iterator** class, that allows us to wrap Rust iterators into Python iterators.

```

for _ in 0..npoints {
    points.push([
        between.sample(&mut range),
        between.sample(&mut range),
        between.sample(&mut range),
    ])
}

// Setup an MPI environment
let universe: Universe = mpi::initialize().unwrap();
let world: UserCommunicator = universe.world();
let comm = world.duplicate();

// Unbalanced
let unbalanced: DistributedTree = DistributedTree::new(&points, false, &comm)

// Balanced
let balanced_tree: DistributedTree = DistributedTree::new(&points, true, &comm)
}

```

Using MPI from Rust

MPI, via the rsMPI crate, constitutes the most complex dependency in Rusty Tree. The MPI bindings are built on top of C shim library, as well as automatically gener-

ated C/C++ compatible header files using the **cbindgen** crate. The shim library is used to form an equivalence between under-specified identifiers from the MPI standard, that may not be uniform across MPI implementations, and **cbindgen** is used to automatically generate headers for different MPI implementations. This allows rsMPI to expose an interface in Rust which covers most of MPI's core functionality for most MPI implementations, on most operating systems ⁶. However, as it is designed to target multiple platforms like all other Rust crates, its build is fragile and requires extra configuration from users who might wish to deploy on unsupported or novel hardware and software targets. Specifically we have found to problems arise on Cray machines, such as Archer 2, which have custom installations of MPI and do not support the **mpicc** compiler wrapper. However, we note that this is an active area of development, and will be accounted for in a future rsMPI release. As a temporary fix, we currently fork of rsMPI with manual configurations to target the HPC systems we wish to run our software on. Consumers of our software will have to follow similar steps if automatic builds do not work, however we note that in the majority of cases this fragility will not be exposed to users, who can continue to who can continue to consume our library using a single line in their project's `Cargo.toml`.

Scaling Experiments

Figure (2.12a) shows the weak scaling performance of Rusty Tree⁷ for constructing *balanced octrees*. $5e5$ particles, taken from a random uniform distribution in a unit cube are placed on each processor, up to a total of 16 million particles. The parameter n_{crit} is set to 200. We compare the performance of Rusty Tree with Dendro, a C++ library similarly built on algorithms presented in [60]. Dendro has a limited API in comparison to Rusty tree, with no Python interface, and no ability to query tree nodes for the points they contain and is therefore not usable for developing fast algorithms. However it does act as a showcase of the algorithms presented in [60], allowing for a direct comparison to our own implementation. We observe that both softwares have a similar scaling behaviour for the tested problem sizes, however Rusty Tree is more expensive by a has a fixed constant factor, both in Python and Rust. We observe that this is due to our provision of hashmaps (see fig. (2.11)), that allow users to query tree nodes for the points they contain and vice versa, a functionality not available in Dendro. An experiment confirming this is shown in figure (2.12d).

We observe that Rusty Tree, when called from Python, has a significantly more expensive memory overhead in comparison to the Rust library, displayed in figure (2.12b). Whereas the Rust library, and Dendro, have similar memory profiles. The cause for excess memory overhead in Python is due to the need to assign static lifetimes to all Rust pointers passed to Python, for example for Morton keys and point data, which would be optimised out of a corresponding Rust program at compile time if it was not being used downstream in a program. We observe that the runtime cost of using Rusty Tree from Python is moderate, at between 5 and 25 % of total runtime in Rust, for the problem sizes tested.

Figure (2.12c) is a plot of the mean difference, of the load on each processor with respect to the mean load for a given geometry, as we scale a problem weakly. The

⁶This is an area of active development, and a list of currently supported features can be found here <https://github.com/rsmpi/rsmpi/blob/main/README.md>

⁷Experiments were taken on an AMD Ryzen Threadripper 3970X 32-Core processor

geometries we choose are the surface of a ‘wiggly torus’, as shown in figure (1.2), and the interior of a unit cube, where we sample $5e5$ points randomly from each domain at each processor. The parameter n_{crit} is again set to 200. Comparing the results from both geometries, we observe a qualitatively weak dependence between the problem size and an increasing load imbalance. However, the overlapping nature of the error bars in our experiments, and the fact that an experiment with uniform random data shows similar load imbalance to a significantly less uniform data from the surface of a wiggly torus, make it difficult to conclude a correlation without further investigation. We will need to supplement this investigation with results for larger experiment sizes on large scale HPC systems.

Conclusion

Rusty Tree has high-level Python interfaces, allowing it to be inserted into existing research projects with minimal configuration with a moderate performance hit. It is designed around data and is simple to read, and importantly can be built for a variety of software and hardware environments. It’s a demonstration of Rust as a viable alternative to development in Rust as opposed to C++ or Fortran. We believe that this will encourage the widespread adoption of our library, and we hope to publicise this software, as well as our user experience with Rust for high performance scientific computing, in an upcoming paper.




Figure 2.12: There are $5e5$ points per processor for the weak scaling experiments in (a), with an $n_{crit} = 200$. Figure (b) shows the peak memory usage of Dendro compared to the Python and Rust interfaces of Rusty Tree. We note that the Python library has a significant overhead. Figure (c) shows the difference in load from the global mean load, as mean over runs, in an experiment that is weakly scaled with the same parameters as in figure (a). Figure (d) shows the breakdown in runtime between the different components of the algorithms for computing balanced and unbalanced trees, experiments are shown for a tree with $2e6$ uniformly randomly distributed points, distributed evenly across 4 processors. We see that the dominant factor in runtime is the creation of the hashmaps.

Looking Ahead

In this chapter we return to the motivating mathematical setting of this thesis, which is the numerical solution of elliptic PDEs using a boundary integral formulation. We begin in section 3.1 by providing an overview of fast direct solvers, which constitute a fast-matrix inversion - in contrast to the fast matrix-vector product of the FMM. This section proceeds to describe ongoing work on the formulation of a new fast direct solver for the scattering of low-frequency described by the Helmholtz equation, in collaboration with researchers at the Flatiron Institute. In section 3.3 we briefly discuss the gap in the literature in the understanding trade-offs between approaches for field translation, which as we saw in chapter 1, are critical for a performant implementations of fast algorithms. We conclude in section 3.4 with a brief summary of our target application, the application of our fast solver infrastructure to solve the scattering of electromagnetic waves which are described by Maxwell's equations.

3.1 Fast Direct Solvers on Distributed Memory Systems

- Introduce the logic behind fast direct solvers via a short literature survey of the most popular methods.
 - Introduce RS-S and skeletonization based approaches, why these are good (proxy compression, can re-use octree data structure, work with moderate frequency oscillatory problems, straightforward to parallelize)
 - Introduce current state of the art work with Manas on proxy compression for Helmholtz problems.
 - Conclude with future plans for fast direct solver using our Galerkin discretized BIE.

Proxy compression is necessary in order to achieve the linear complexity bound of the fast-direct solver powered by RS-S. The idea rests on the principle of representing the far-field particles of a given box B , which may contain $O(N)$ particles, with a set of 'proxy points' contained on a proxy surface that encloses B . This surface is often chosen to be a sphere. By choosing $O(1)$ proxy points, without getting into the details yet of how exactly they are sampled, we are able to obtain the linear complexity we desire.

For a given box B , a proxy surface D and its boundary γ are chosen such that $B \subset D$. The far-field points of B , \mathcal{F} is partitioned such that $\mathcal{F} = \mathcal{Q} \cup \mathcal{P}$, where \mathcal{Q} contains $O(1)$ points. Γ is the boundary of the entire scatterer, and $\tau = \Gamma \cap B$ is the portion of the scatterer boundary contained in B . The situation is sketched in figure (??) in 2D.

We can choose to represent our solution due to the charge in B in \mathcal{F} however we

wish. However, our choice will lead to different matrices that we must compress.

Generally, we'll end up with a solution matrix of the form $A_{\mathcal{F}B}$ that maps between B and points in its far-field that can be split up as,

$$v_{\mathcal{F}} = A_{\mathcal{F}B}\psi_B \quad (3.1)$$

$$= B_{\mathcal{F}\gamma}C_{\gamma B}\psi_B \quad (3.2)$$

the subscripts indicate the domains these operators map between. We desire a split like this, as the far field interaction of B can be compressed into something involving $C_{\gamma B}$.

To see this, consider the fact that $A_{\mathcal{F}B}$ can be written as,

$$A_{\mathcal{F}B} = \begin{bmatrix} A_{\mathcal{Q}B} \\ A_{\mathcal{P}B} \end{bmatrix} = \begin{bmatrix} I & 0 \\ 0 & B_{\mathcal{P}\gamma} \end{bmatrix} \begin{bmatrix} A_{\mathcal{Q}B} \\ C_{\gamma B} \end{bmatrix} \quad (3.3)$$

Our RS-S algorithm relies on a compression of this matrix, however a direct compression of $A_{\mathcal{F}B}$ is too expensive, as there are typically $O(N)$ points in \mathcal{F} for a box B , therefore assembly of this matrix for all boxes will result in an algorithm of $O(N^2)$ complexity. However if we can find a decomposition like above, we can apply an interpolative decomposition to the right column in (3.3) which has dimensions $O(1) \times O(n_{\gamma})$ by construction where n_{γ} is the number of proxy points. To prove that this allows us to reconstruct the full matrix after compression. Consider an ID that gives us,

$$\begin{bmatrix} A_{\mathcal{Q}B} \\ C_{\gamma B} \end{bmatrix} = \begin{bmatrix} A_{\mathcal{Q}S} \\ C_{\gamma S} \end{bmatrix} \begin{bmatrix} T_{SR} & 1 \end{bmatrix} \quad (3.4)$$

Where S and R are the skeleton and redundant points respectively. Plugging back into our expression (3.3),

$$A_{\mathcal{F}B} = \begin{bmatrix} I & 0 \\ 0 & B_{\mathcal{P}\gamma} \end{bmatrix} \begin{bmatrix} A_{\mathcal{Q}S} \\ C_{\gamma S} \end{bmatrix} \begin{bmatrix} T_{SR} & 1 \end{bmatrix} \quad (3.5)$$

$$= \begin{bmatrix} A_{\mathcal{Q}S} \\ B_{\mathcal{P}\gamma}C_{\gamma S} \end{bmatrix} \begin{bmatrix} T_{SR} & 1 \end{bmatrix} \quad (3.6)$$

$$= A_{\mathcal{F}S} \begin{bmatrix} T_{SR} & 1 \end{bmatrix} \quad (3.7)$$

Therefore, we see that we can get away with a cheap ID to reconstruct the far-field operator, involving the proxy points rather than the full far field of B .

3.1.1 \mathcal{T}

Outgoing Skeletonization

A double-layer potential, due to some unknown density ψ , supported on τ ,

$$v(x) = \int_{\Gamma \cap B} \frac{\partial \Phi(x, y)}{\partial n(y)} \psi(y) ds(y) := \mathcal{D}\psi, \quad x \in \mathbb{R}^m \setminus \tau \quad (3.8)$$

solves the Helmholtz equation everywhere it's valid. Here, $\Phi(x, y)$ is the fundamental solution of the Helmholtz equation. However, its normal derivative evaluated at the target points, which we'll need for deriving boundary integral equations for Maxwell problems, does not,

$$\frac{\partial v}{\partial n(x)} = \frac{\partial}{\partial n(x)} \int_{\Gamma \cap B} \frac{\partial \Phi(x, y)}{\partial n(y)} \psi(y) ds(y) := \mathcal{T}\psi, \quad x \in \Gamma \cap \mathcal{F} \quad (3.9)$$

it's only valid at far-field points, $\Gamma \cap \mathcal{F}$. However, we can separate out the normal part of the derivative,

$$\frac{\partial v}{\partial n(x)} = n(x) \cdot \nabla_x \int_{\Gamma \cap B} \frac{\partial \Phi(x, y)}{\partial n(y)} \psi(y) ds(y) := n \cdot w \quad (3.10)$$

The function

$$w(x) = \nabla_x \int_{\Gamma \cap B} \frac{\partial \Phi(x, y)}{\partial n(y)} \psi(y) ds(y) := \nabla_x \mathcal{D}\psi \quad (3.11)$$

Does satisfy our PDE, everywhere, and we'll exploit this fact in a moment. As an aside, we can see that this is true by considering a double layer potential v that is smooth enough to admit,

$$(\Delta + k^2)w = (\Delta + k^2)\nabla_x v = \nabla_x(\Delta + k^2)v = 0 \quad (3.12)$$

where the last equality follows as v satisfies the Helmholtz equation. Therefore w is a solution of the Helmholtz equation. Note that w has three components.

In order to find our $C_{\gamma B}$ with this representation, we need to set up an 'associated boundary value problem' for each component of w . The choice of boundary value problem we choose is free, as we only rely on the existence of its solution.

Consider an associated boundary value problem for just a single component of \tilde{w} that satisfies,

$$(\Delta + k^2)\tilde{w} = 0, \quad x \in \mathbb{R}^m \setminus D \quad (3.13)$$

$$\tilde{w} = w_1(x) \quad (3.14)$$

$$\text{A radiation condition at } \infty \quad (3.15)$$

A combined field representation might be nice, as we know it has good properties,

$$\tilde{w} = (\mathcal{D} - ik\mathcal{S})_{\mathcal{F}\gamma} \mu \quad (3.16)$$

where μ is some unknown density supported on the proxy surface γ . Forming the boundary integral equation, and plugging back into the representation for \tilde{w} ,

$$\tilde{w} = (\mathcal{D} - ik\mathcal{S})_{\mathcal{F}\gamma} \left(\frac{1}{2}\mathcal{I} + \mathcal{D} - ik\mathcal{S} \right)_{\gamma\gamma}^{-1} w_1 \quad (3.17)$$

$$= (\mathcal{D} - ik\mathcal{S})_{\mathcal{F}\gamma} \left(\frac{1}{2}\mathcal{I} + \mathcal{D} - ik\mathcal{S} \right)_{\gamma\gamma}^{-1} \nabla_1 \mathcal{D}_{\gamma B} \psi_\gamma \quad (3.18)$$

$$\equiv B_{\mathcal{F}\gamma} C_{\gamma B} \psi_\gamma \quad (3.19)$$

where we identify,

$$C_{\gamma B} = \nabla_1 \mathcal{D}_{\gamma B} \quad (3.20)$$

This is the matrix we will attempt to compress. Similar analysis follows for the other two components of $w(x)$. Meaning that we end up having to compress $[\nabla_1 \mathcal{D}_{\gamma B}, \nabla_2 \mathcal{D}_{\gamma B}, \nabla_3 \mathcal{D}_{\gamma B}]$ for the outgoing problem.

We see that $B_{\mathcal{F}\gamma}$ is never explicitly formed, we just require its existence. When we calculate an approximation of $A_{\mathcal{F}B}$ using (3.5), we only need to know the ID of the $C_{\gamma B}$.

Incoming Skeletonization

For the incoming skeletonization, were again we're considering the same representation with a hypersingular operator, we observe that we're just looking for,

$$\left[\frac{\partial v}{\partial n(x)} \right]_{\mathcal{F}B}^T \quad (3.21)$$

with the formation of an associated boundary integral equation taking place in much the same way as for the outgoing problem. However, the hypersingular operator is self-adjoint, therefore it leads to the same expressions for $C_{\gamma B}$.

3.1.2 \mathcal{K}'

Outgoing Skeletonization

If we choose to represent our potential with a single-layer potential,

$$u(x) = \int_{\Gamma \cap B} \Phi(x, y) \phi(y) ds(y) := \mathcal{S}\phi, \quad x \in \mathbb{R}^m \setminus \tau \quad (3.22)$$

and seek a boundary integral equation in terms of its normal derivative at the targets,

$$w(x) = \int_{\Gamma \cap B} \frac{\partial \Phi(x, y)}{\partial n(x)} \phi(y) ds(y) := \mathcal{K}'\phi, \quad x \in \Gamma \cap \mathcal{F} \quad (3.23)$$

We observe the same problem as in the \mathcal{T} case, where this expression is not a general solution of our PDE. We can similarly separate out the normal component and write,

$$\tilde{w}(x) := \int_{\Gamma \cap B} \nabla_x \Phi(x, y) \phi(y) ds(y), \quad x \in \mathbb{R}^m \setminus \tau \quad (3.24)$$

Using the previous analysis for \mathcal{T} , we immediately recognise that the components we must compress are $C_{\gamma B} = \nabla_1 \mathcal{S}_{\gamma B}$, giving us $[\nabla_1 \mathcal{S}_{\gamma B}, \nabla_2 \mathcal{S}_{\gamma B}, \nabla_3 \mathcal{S}_{\gamma B}]$ to compress in total for the outgoing problem.

Incoming Skeletonization

Noticing that,

$$\left[\frac{\partial u}{\partial n(x)} \right]_{\mathcal{F}B}^T = \int_{\Gamma \cap B} \frac{\partial \Phi(x, y)}{\partial n(y)} \phi(y) ds(y) = \mathcal{D}_{\gamma B} \phi \quad (3.25)$$

already satisfies our PDE without any further work, we can save a lot of work, and simply use it as our Dirichlet data in the associated boundary value problem. The matrix to compress being $C_{\gamma B} = \mathcal{D}_{\gamma B}$.

3.1.3 \mathcal{S}

Outgoing Skeletonization

We now choose to represent our scattered solution with a single-layer operator,

$$u(x) = \int_{\Gamma \cap B} \Phi(x, y) \phi(y) ds(y) := \mathcal{S} \phi, \quad x \in \mathbb{R}^m \setminus \tau \quad (3.26)$$

This satisfies the underlying PDE everywhere. We can now set up an associated exterior boundary value problem as before, and use our single-layer potential as Dirichlet boundary data.

$$(\Delta + k^2)w = 0, \quad x \in \mathbb{R}^m \setminus D \quad (3.27)$$

$$w = \mathcal{S} \phi, \quad \text{on } \gamma \quad (3.28)$$

$$\text{Radiation condition at } \infty \quad (3.29)$$

where ϕ is some unknown density supported on τ . As before, we can form a boundary integral equation for this associated problem, and solve, recognising that the matrix to compress $C_{\gamma B} = \mathcal{S}_{\gamma B}$

Incoming Skeletonization

The single-layer operator is self-adjoint, leading to the same operator to compress. Spelling this out, consider the associated interior boundary value problem,

$$(\Delta + k^2)w = 0 \quad \text{in } D \quad (3.30)$$

$$w = \mathcal{S} \phi \quad \text{on } \gamma \quad (3.31)$$

The solution of an interior Helmholtz scattering problem may not be unique, but this doesn't matter for our purposes. Proxy compression doesn't require uniqueness, only existence. Let's seek a solution in the form of a combined-layer potential,

$$w(x) = (\mathcal{D}_{\gamma} - ik\mathcal{S}_{\gamma})[\phi](x) \quad (3.32)$$

where the subscripts make it clear that the density is supported on γ . Forming the boundary integral equation,

$$\left(-\frac{\mathcal{I}}{2} + \mathcal{D}_\gamma - ik\mathcal{S}_\gamma\right)[\phi](x) = \mathcal{S}_{\mathcal{F}\gamma}[\phi](x) \quad (3.33)$$

Solving with the representation gives,

$$w = \mathcal{S}_{B\gamma}\left(-\frac{\mathcal{I}}{2} + \mathcal{D}_\gamma - ik\mathcal{S}_\gamma\right)^{-1}\mathcal{S}_{\mathcal{F}\gamma}[\phi](x) = C_{B\gamma}B_{\gamma\mathcal{F}} \quad (3.34)$$

where we recognize the matrix to compress as $C_{B\gamma} = \mathcal{S}_{B\gamma}$ in our proxy framework.

3.1.4 \mathcal{D}

Outgoing Skeletonization

A double-layer potential, due to some unknown density ψ , supported on τ ,

$$v(x) = \int_{\Gamma \cap B} \frac{\partial \Phi(x, y)}{\partial n(y)} \psi(y) ds(y) := \mathcal{D}\psi, \quad x \in \mathbb{R}^m \setminus \tau \quad (3.35)$$

solves the Helmholtz equation everywhere it's valid. Therefore it can be used as Dirichlet data for the associated boundary value problem for the outgoing skeletonization. Applying similar analysis to above, we identify the kernel to compress as $C_{\gamma B} = \mathcal{D}_{\gamma B}$.

Incoming Skeletonization

We notice that the transpose of the double layer operator is,

$$[u]_{\mathcal{F}B}^T(x) = \int_{\Gamma \cap B} \frac{\partial \Phi(x, y)}{\partial n(x)} \phi(y) ds(y) = \mathcal{K}'_{\gamma B} \phi \quad (3.36)$$

This in general does not satisfy our PDE everywhere, we again separate out the normal component, and as before, recognise that the components to compress are $[\nabla_1 \mathcal{S}_{\gamma B}, \nabla_2 \mathcal{S}_{\gamma B}, \nabla_2 \mathcal{S}_{\gamma B}]$.

3.2 Example Problems

Let's now apply our fast direct solver framework, with proxy compression to some example problems. We begin with acoustic sound-hard scattering, which is a didactic example.

3.2.1 Acoustic Sound Hard Scattering

Consider a scattered field u^s , that bounces off an object Ω and satisfies the Helmholtz equation in the exterior,

$$(\Delta + k^2)u^s = 0, \quad \text{in } \mathbb{R}^3 \setminus \Omega \quad (3.37)$$

The ‘sound hard’ boundary condition on the surface Γ is

$$\frac{\partial u^s}{\partial n} = \frac{\partial u^i}{\partial n}, \quad \text{in } \Gamma \quad (3.38)$$

where u^i is the incident wave. Using the analysis in [Bruno2012], we write down a ‘regularised’ representation formula for our solution. This regularisation can be shown to have better spectral properties.

$$u^s = (\mathcal{K}_k \circ \mathcal{S}_K - i\eta \mathcal{S}_k) \quad (3.39)$$

where k and K are complex wave numbers, that may not be the same. We can take the trace of this representation, and its normal derivative at the targets, and find a boundary integral equation for the exterior problem,

$$\left(\frac{i\eta}{2}\mathcal{I} - i\eta\mathcal{K}'_k + \mathcal{T}_k \circ \mathcal{S}_K\right)\mu = g \quad (3.40)$$

using the Calderón identity,

$$\mathcal{T}_k \circ \mathcal{S}_k = -\frac{1}{4}\mathcal{I} + (K'_k)^2 \quad (3.41)$$

which is true for any k , we arrive at a boundary integral equation,

$$\left(i\eta\left(\frac{1}{2}\mathcal{I} - \mathcal{K}'_k\right) - \frac{1}{4}I + (K'_k)^2\right)\mu = g \quad (3.42)$$

by defining $\theta := \mathcal{K}'_k\mu$, we can write the boundary integral equation as a system,

$$\begin{pmatrix} \left(\frac{i\eta}{2} - \frac{1}{4}\right)\mathcal{I} - i\eta\mathcal{K}'_k & \mathcal{K}'_k \\ \mathcal{K}'_k & -\mathcal{I} \end{pmatrix} \begin{pmatrix} \mu \\ \theta \end{pmatrix} = \begin{pmatrix} g \\ 0 \end{pmatrix} \quad (3.43)$$

We can then place this system into our fast direct solver framework. Despite not knowing how to compress the system matrix altogether, we do know how to compress each block, as they each correspond to displacements of \mathcal{K}'_k .

Consider writing out our block matrix as,

$$\begin{pmatrix} A & B \\ C & D \end{pmatrix} \begin{pmatrix} \mu \\ \theta \end{pmatrix} = \begin{pmatrix} g \\ 0 \end{pmatrix} \quad (3.44)$$

and re-writing as,

$$\begin{pmatrix} \begin{pmatrix} A_{11} & B_{11} \\ C_{11} & D_{11} \end{pmatrix} & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \begin{pmatrix} A_{NN} & B_{NN} \\ C_{NN} & D_{NN} \end{pmatrix} \end{pmatrix} \begin{pmatrix} \mu_1 \\ \theta_1 \\ \cdot \\ \mu_n \\ \theta_n \end{pmatrix} = \begin{pmatrix} g_1 \\ 0 \\ \cdot \\ g_N \\ 0 \end{pmatrix} \quad (3.45)$$

This system matrix remains numerically low-rank, and therefore can fit into our FDS framework. Indeed, we now have to compress a system of matrix operators in order to capture its row space via the proxy trick,

3.2.2 Acoustic Transmission

The sound-hard problem is not particularly realistic, the transmission problem accounts for different material properties in the scatterer.

This time, we get two Helmholtz equations, for the interior of the scatterer Ω , and unbounded exterior.

$$(\Delta + k^2)u(x) = 0, \quad x \in \mathbb{R}^3 \setminus \Omega \quad (3.46)$$

$$(\Delta + k_0^2)u_0(x) = 0, \quad x \in \Omega \quad (3.47)$$

The boundary conditions are,

$$\mu u - \mu_0 u_0 = f, \quad \text{on } \partial\Omega \quad (3.48)$$

$$\frac{\partial u}{\partial \nu} - \frac{\partial u_0}{\partial \nu} = g, \quad \text{on } \partial\Omega \quad (3.49)$$

Consider the exterior scattering problem. We seek a solution in terms of combined single and double layer surface potentials,

$$u = \mathcal{K}\psi + \mu\mathcal{S}\phi, \quad \text{in } \mathbb{R}^3 \setminus \Omega \quad (3.50)$$

$$u_0 = \mathcal{K}_0\psi + \mu_0\mathcal{S}_0\phi, \quad \text{in } \Omega \quad (3.51)$$

$$(3.52)$$

Where the subscript, 0, denotes that the fundamental solution has been calculated with a different values of wave number, k_0 .

Using the jump relations, we can form a linear system from the resulting boundary integral equations. First, we obtain two simultaneous boundary integral equations,

$$(\mu + \mu_0)\psi + (\mu\mathcal{K} - \mu_0\mathcal{K}_0)\psi + (\mu^2\mathcal{S} - \mu_0^2\mathcal{S}_0)\phi = 2f \quad (3.53)$$

$$(\mu + \mu_0)\phi - (\mathcal{T} - \mathcal{T}_0)\psi - (\mu\mathcal{K}' - \mu_0\mathcal{K}'_0)\phi = -2g \quad (3.54)$$

Defining a system matrix,

$$\mathbf{A}' := \begin{pmatrix} -(\mu\mathcal{K} - \mu_0\mathcal{K}_0) & -(\mu^2\mathcal{S} - \mu_0^2\mathcal{S}_0) \\ \mathcal{T} - \mathcal{T}_0 & (\mu\mathcal{K}' - \mu_0\mathcal{K}'_0) \end{pmatrix} \quad (3.55)$$

and we can write out the linear system to be solved as,

$$(\mu + \mu_0)\xi - \mathbf{A}'\xi = 2h \quad (3.56)$$

where,

$$\xi = \begin{pmatrix} \psi \\ \phi \end{pmatrix} \quad \text{and} \quad h = \begin{pmatrix} f \\ -g \end{pmatrix} \quad (3.57)$$

This is of the form $I + \text{compact}$ and can therefore be inserted into Fredholm theory. In practice, it may be easier to work with the system matrix written as,

$$\mathbf{A} := \begin{pmatrix} (\mu + \mu_0) + (\mu\mathcal{K} - \mu_0\mathcal{K}_0) & (\mu^2\mathcal{S} - \mu_0^2\mathcal{S}_0) \\ -(\mathcal{T} - \mathcal{T}_0) & (\mu + \mu_0) - (\mu\mathcal{K}' - \mu_0\mathcal{K}'_0) \end{pmatrix} \quad (3.58)$$

We're now in a position to start considering how we may apply proxy compression to this block matrix. We already know how to compress all the blocks separately, so we can now proceed as we did for the sound-hard problem.

3.3 Optimal Translation Operators for Fast Algorithms

As mentioned in chapter 1, the FMM and other fast algorithms can be optimised by their field translations, some of which were discussed in that chapter. As it stands there is a gap in the literature that quantitatively compares and contrasts the relative merits of different approaches to field translation. We summarise some of the popular strategies in more detail below.

- Translation operators, what are they, and what are the different approaches currently used.
- What are the trade-offs of different approaches?
- Can I write some quick software for the quick comparison of translation operators - maybe in Python, on top of RustyTree? This would allow me to get some graphs to compare between approaches. If this is too much work, I will have to just compare the approaches in words.

As translation operators are agnostic to the actual methods used to approximate fields, whether by a multipole expansions, or a black box method, a trait based library would provide the perfect interface for users to configure different operators to their particular geometries. We provide an example of how this interface could look in figure [TODO: translation library trait syntax]. A future vector of research would be to complete this study, and release this library, **rusty-field**, as a component software of our fast solver infrastructure.

3.4 Target Application: Maxwell Scattering

- Very brief summary of the Maxwell scattering problem, how we will form the BIE, the representation formulae we'll use, and how the integral operator will be discretised.
- Overview of what kind of problems this would help us solve?

Conclusion

In this subsidiary thesis we've presented progress on the development of a new software infrastructure for fast algorithms. We've documented recent software outputs towards this goal, as well as plans for future mathematical and software developments. The main outputs being an investigation into programming languages and environments most suitable for scientific computing, a parallel load balanced octree library designed for high-performance, as well as a mathematical investigation of fast direct solvers for Helmholtz problems. The immediate next steps of this project will be to publish our recent software results on octrees, and mathematical developments on fast direct solvers in appropriate scientific journals. In the medium term, we plan to complete the implementation of the parallel FMM software, and the apply this to large scale boundary integral equation problems. This would entail the completion of the majority of our fast solver infrastructure. The final stage of this project will be the extension of our infrastructure to a parallel fast direct solver, and the simulation of large scale electromagnetic scattering problems.

Appendix

A.1 Deriving Local Expansion Coefficients from Multipole Expansion in \mathbb{R}^2

A.2 Complexity of Tree Construction

A.3 HykSort

HykSort algorithm

Parallel Selection algorithm

Hyksort vs Sample sort complexity

Bibliography

- [1] Sivaram Ambikasaran. *Fast algorithms for dense numerical linear algebra and applications*. Stanford University, 2013.
- [2] Sivaram Ambikasaran and Eric Darve. “An $O(N \log N)$ Fast Direct Solver for Partial Hierarchically Semi-Separable Matrices”. In: *Journal of Scientific Computing* 57.3 (2013), pp. 477–501.
- [3] Sivaram Ambikasaran et al. “Large-scale stochastic linear inversion using hierarchical matrices”. In: *Computational Geosciences* 17.6 (2013), pp. 913–927.
- [4] *B2: makes it easy to build C++ projects, everywhere*. 2022. URL: <https://github.com/boostorg/build>.
- [5] *Bazel - a fast, scalable, multi-language and extensible build system*. Version 5.3.2. 2022. URL: <https://github.com/bazelbuild/bazel>.
- [6] *BLAS and LAPACK for Rust*. 2022. URL: <https://github.com/blas-lapack-rs>.
- [7] Steffen Börm, Lars Grasedyck, and Wolfgang Hackbusch. “Introduction to hierarchical matrices with applications”. In: *Engineering analysis with boundary elements* 27.5 (2003), pp. 405–422.
- [8] Steffen Börm and Wolfgang Hackbusch. “A short overview of H2-matrices”. In: *Proceedings in applied mathematics and mechanics* 2.1 (2003), pp. 33–36.
- [9] Carsten Burstedde, Lucas C. Wilcox, and Omar Ghattas. “p4est: Scalable Algorithms for Parallel Adaptive Mesh Refinement on Forests of Octrees”. In: *SIAM Journal on Scientific Computing* 33.3 (2011), pp. 1103–1133. DOI: 10.1137/100791634.
- [10] Stéphanie Chaillat, Marc Bonnet, and Jean-François Semblat. “A multi-level fast multipole BEM for 3-D elastodynamics in the frequency domain”. In: *Computer Methods in Applied Mechanics and Engineering* 197.49-50 (2008), pp. 4233–4249.
- [11] Shiv Chandrasekaran et al. “A fast solver for HSS representations via sparse matrices”. In: *SIAM Journal on Matrix Analysis and Applications* 29.1 (2007), pp. 67–81.
- [12] *Conan - The open-source C/C++ package manager*. Version 1.53.0. 2022. URL: <https://github.com/conan-io/conan>.
- [13] Lisandro Dalcin and Yao-Lung L Fang. “mpi4py: Status update after 12 years of development”. In: *Computing in Science & Engineering* 23.4 (2021), pp. 47–54.

- [14] Eric Darve and Pascal Havé. “A fast multipole method for Maxwell equations stable at all frequencies”. In: *Philosophical Transactions of the Royal Society of London. Series A: Mathematical, Physical and Engineering Sciences* 362.1816 (2004), pp. 603–628.
- [15] *Dendro - Distributed adaptive octree construction*. Version 2.0. 2020. URL: <https://github.com/paralab/Dendro-5.01>.
- [16] Jack Dongarra et al. “With extreme computing, the rules have changed”. In: *Computing in Science & Engineering* 19.3 (2017), pp. 52–62.
- [17] Björn Engquist and Lexing Ying. “Fast directional multilevel algorithms for oscillatory kernels”. In: *SIAM Journal on Scientific Computing* 29.4 (2007), pp. 1710–1737.
- [18] William Fong and Eric Darve. “The black-box fast multipole method”. In: *Journal of Computational Physics* 228.23 (2009), pp. 8712–8725.
- [19] *Fortran Package Manager (fpm)*. Version 0.7.0. 2022. URL: <https://github.com/fortran-lang/fpm>.
- [20] Yuhong Fu and Gregory J Rodin. “Fast solution method for three-dimensional Stokesian many-particle problems”. In: *Communications in Numerical Methods in Engineering* 16.2 (2000), pp. 145–149.
- [21] Yuhong Fu et al. “A fast solution method for three-dimensional many-particle problems of linear elasticity”. In: *International Journal for Numerical Methods in Engineering* 42.7 (1998), pp. 1215–1229.
- [22] Amir Gholami et al. “AccFFT: A library for distributed-memory FFT on CPU and GPU architectures”. In: *arXiv preprint arXiv:1506.07933* (2015).
- [23] Amir Gholami et al. “FFT, FMM, or multigrid? A comparative study of state-of-the-art Poisson solvers for uniform and nonuniform grids in the unit cube”. In: *SIAM Journal on Scientific Computing* 38.3 (2016), pp. C280–C306.
- [24] Pieter Ghysels et al. “STRUMPACK: Scalable Preconditioning using Low-Rank Approximations and Random Sampling”. In: ().
- [25] Zydrunas Gimbutas and Vladimir Rokhlin. “A generalized fast multipole method for nonoscillatory kernels”. In: *SIAM Journal on Scientific Computing* 24.3 (2003), pp. 796–817.
- [26] Leslie Greengard and Vladimir Rokhlin. “A fast algorithm for particle simulations”. In: *Journal of computational physics* 73.2 (1987), pp. 325–348.
- [27] Leslie F Greengard and Jingfang Huang. “A new version of the fast multipole method for screened Coulomb interactions in three dimensions”. In: *Journal of Computational Physics* 180.2 (2002), pp. 642–658.
- [28] Wolfgang Hackbusch. “A sparse matrix arithmetic based on H-matrices. part i: Introduction to H-matrices”. In: *Computing* 62.2 (1999), pp. 89–108.
- [29] Sijia Hao, Per-Gunnar Martinsson, and Patrick Young. “An efficient and highly accurate solver for multi-body acoustic scattering problems involving rotationally symmetric scatterers”. In: *Computers & Mathematics with Applications* 69.4 (2015), pp. 304–318.
- [30] Kenneth L Ho. “FLAM: Fast linear algebra in MATLAB-Algorithms for hierarchical matrices”. In: *Journal of Open Source Software* 5.51 (2020), p. 1906.

- [31] Srinath Kailasa et al. “PyExaFMM: an exercise in designing high-performance softawre with Python and Numba”. In: *To Appear in Computing in Science and Engineering* 24.4 (2022).
- [32] Steve Klabnik and Carol Nichols. *The Rust Programming Language (Covers Rust 2018)*. No Starch Press, 2019.
- [33] Siu Kwan Lam, Antoine Pitrou, and Stanley Seibert. “Numba: A llvm-based python jit compiler”. In: *Proceedings of the Second Workshop on the LLVM Compiler Infrastructure in HPC*. 2015, pp. 1–6.
- [34] Chris Lattner and Vikram Adve. “LLVM: A compilation framework for lifelong program analysis & transformation”. In: *International Symposium on Code Generation and Optimization, 2004. CGO 2004*. IEEE. 2004, pp. 75–86.
- [35] Dongryeol Lee, Richard Vuduc, and Alexander G Gray. “A distributed kernel summation framework for general-dimension machine learning”. In: *Proceedings of the 2012 SIAM International Conference on Data Mining*. SIAM. 2012, pp. 391–402.
- [36] Denis Zorin Lexing Ying George Biros. “A kernel-independent adaptive fast multipole algorithm in two and three dimensions”. In: *Journal of Computational Physics* 196.2 (2004), pp. 591–626. DOI: <http://dx.doi.org/10.1016/j.jcp.2003.11.021>.
- [37] Judith Yue Li et al. “A Kalman filter powered by-matrices for quasi-continuous data assimilation problems”. In: *Water Resources Research* 50.5 (2014), pp. 3734–3749.
- [38] Anton Malakhov. “Composable Multi-Threading for Python Libraries”. In: *Proceedings of the 15th Python in Science Conference* (2016), pp. 15–19. DOI: 10.25080/majora-629e541a-002. URL: <https://youtu.be/kfQcWez2URE>.
- [39] Dhairya Malhotra and George Biros. “PVFMM: A parallel kernel independent FMM for particle and volume potentials”. In: *Communications in Computational Physics* 18.3 (2015), pp. 808–830.
- [40] Per-Gunnar Martinsson and Vladimir Rokhlin. “An accelerated kernel-independent fast multipole method in one dimension”. In: *SIAM Journal on Scientific Computing* 29.3 (2007), pp. 1160–1178.
- [41] *Maturin*. Version 0.13.0. 2022. URL: <https://github.com/PyO3/maturin>.
- [42] Matthias Messner et al. “Optimized M2L kernels for the Chebyshev interpolation based fast multipole method”. In: *arXiv preprint arXiv:1210.7292* (2012).
- [43] *Microsoft: 70 percent of all security bugs are memory safety issues*. <https://www.zdnet.com/article/microsoft-70-percent-of-all-security-bugs-are-memory-safety-issues/>. Accessed: 2022-10-31.
- [44] Victor Minden. *strong-skel*. Version 0.1.0. Dec. 15, 2018. URL: <https://github.com/victorminden/strongskel>.
- [45] Victor Minden et al. “A recursive skeletonization factorization based on strong admissibility”. In: *Multiscale Modeling & Simulation* 15.2 (2017), pp. 768–796.
- [46] *ndarray: an N-dimensional array with array views, multidimensional slicing, and efficient operations*. Version 0.15.6. 2022. URL: <https://github.com/rust-ndarray/ndarray>.

- [47] *PyO3: Rust bindings for the Python interpreter*. Version 0.17.3. Nov. 1, 2022. URL: <https://github.com/PyO3/pyo3>.
- [48] Manas Rachh et al. *FMM3D*. Version 1.0.1. June 29, 2022. URL: <https://github.com/flatironinstitute/FMM3D>.
- [49] Manas Rachh et al. *FMM3D*. Version 1.0.1. June 29, 2022. URL: <https://github.com/flatironinstitute/FMM3D>.
- [50] Manas Rachh et al. *FMM3DBIE*. Version 0.1.0. Oct. 5, 2022. URL: <https://github.com/fastalgorithms/fmm3dbie>.
- [51] Abtin Rahimian et al. “Petascale direct numerical simulation of blood flow on 200k cores and heterogeneous architectures”. In: *SC’10: Proceedings of the 2010 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis*. IEEE. 2010, pp. 1–11.
- [52] *Rust-CPython: Rust-Python bindings*. Version 0.7.1. Oct. 25, 2022. URL: <https://github.com/dgrunwald/rust-cpython>.
- [53] *SCons - a software construction tool*. Version 4.4.0. 2022. URL: <https://github.com/SCons/scons>.
- [54] Craig Scott. *Professional CMake: A Practical Guide*. 2018.
- [55] *Spack - A flexible package manager that supports multiple versions, configurations, platforms, and compilers*. Version 0.18.1. 2022. URL: <https://github.com/spack/spack>.
- [56] Benedikt Steinbusch and Andrew Gaspar et al. *RSMPI: MPI bindings for Rust*. Version 0.5.4. 2018. URL: <https://github.com/rsmpi/rsmpi>.
- [57] Josh Stone and Niko Matsakis et. al. *Rayon: A data parallelism library for Rust*. Version 1.5.3. 2022. URL: <https://github.com/rayon-rs/rayon>.
- [58] Hansol Suh and Tobin Isaac. “Evaluation of a minimally synchronous algorithm for 2: 1 octree balance”. In: *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis*. IEEE. 2020, pp. 1–12.
- [59] Hari Sundar, Dhairya Malhotra, and George Biros. “Hyksort: a new variant of hypercube quicksort on distributed memory architectures”. In: *Proceedings of the 27th international ACM conference on international conference on supercomputing*. 2013, pp. 293–302.
- [60] Hari Sundar, Rahul S Sampath, and George Biros. “Bottom-up construction and 2: 1 balance refinement of linear octrees in parallel”. In: *SIAM Journal on Scientific Computing* 30.5 (2008), pp. 2675–2708.
- [61] *The Meson Build System*. Version 0.63.3. 2022. URL: <https://github.com/mesonbuild/meson>.
- [62] Tiankai Tu, David R O’Hallaron, and Omar Ghattas. “Scalable parallel octree meshing for terascale applications”. In: *SC’05: Proceedings of the 2005 ACM/IEEE conference on Supercomputing*. IEEE. 2005, pp. 4–4.
- [63] *VCPKG - C++ Library Manager for Windows, Linux, and MacOS*. Version 2022.10.19. 2022. URL: <https://github.com/microsoft/vcpkg>.
- [64] Bruce Wagar. “Hyperquicksort: A fast sorting algorithm for hypercubes”. In: *Hypercube Multiprocessors* 1987 (1987), pp. 292–299.

- [65] Tingyu Wang, Rio Yokota, and Lorena A Barba. “ExaFMM: a high-performance fast multipole method library with C++ and Python interfaces”. In: *Journal of Open Source Software* 6.61 (2021), p. 3145.
- [66] Tingyu Wang et al. “High-productivity, high-performance workflow for virus-scale electrostatic simulations with Bempp-Exafmm”. In: *arXiv preprint arXiv:2103.01048* (2021).
- [67] William R Wolf and Sanjiva K Lele. “Aeroacoustic integrals accelerated by fast multipole method”. In: *AIAA journal* 49.7 (2011), pp. 1466–1477.
- [68] Rio Yokota, Huda Ibeid, and David Keyes. “Fast multipole method as a matrix-free hierarchical low-rank approximation”. In: *International Workshop on Eigenvalue Problems: Algorithms, Software and Applications in Petascale Computing*. Springer. 2015, pp. 267–286.
- [69] Yokota, Rio and Wang, Tingu and Zhang, Chen Wu and Barba, Lorena A. *ExaFMM*. Version 0.1.0. Oct. 5, 2022. URL: <https://github.com/exafmm/exafmm>.