

Modern Research Software For Fast Multipole Methods

Srinath Kailasa

A thesis submitted in partial fulfillment of the requirements
for the degree of Doctor of Philosophy

Department of Mathematics
University College London

September, 2024

Declaration

I, Srinath Kailasa, confirm that the work presented in this thesis is my own. Where information has been derived from other sources, I confirm that this has been indicated in the thesis.

Acknowledgements

The completion of this thesis simply wouldn't have been possible without the strong prevailing wind of emotional support from my family the regularity of fun with my friends, and the sympathetic and dedicated teachers and colleagues I met at UCL and across the world. Thank you *all* for giving me this opportunity, I'm excited for what the future brings.

నేను ఇది మా అమ్మ కోసం రాశాను, మీరు లేకుండా ఇది జరిగేది కాదు |

UCL Research Paper Declaration Form

Published Manuscripts

1. PyExaFMM: an exercise in designing high-performance software with Python and Numba.
 - a) DOI: 10.1109/MCSE.2023.3258288
 - b) Journal: Computing in Science & Engineering
 - c) Publisher: IEEE
 - d) Date of Publication: Sept-Oct 2022
 - e) Authors: Srinath Kailasa, Tingyu Wang, Lorena A. Barba, Timo Betcke
 - f) Peer Reviewed: Yes
 - g) Copyright Retained: No
 - h) ArXiv: 10.48550/arXiv.2303.08394
 - i) Associated Thesis Chapters: 2
 - j) Statement of Contribution
 - i. Srinath Kailasa was the lead author and responsible for the direction of this research and the preparation of the manuscript.
 - ii. Tingyu Wang offered expert guidance as on fast multipole method implementations, and critical feedback of the manuscript.
 - iii. Lorena A. Barba served as an advisor, offering valuable insights into the structure of the manuscript, suggesting improvements to enhance clarity and impact of the work.
 - iv. Timo Betcke provided significant advisory support, contributing to the design and framing of the research, interpretation of the results and provided critical feedback of the manuscript.

Unpublished Manuscripts

1. M2L Translation Operators for Kernel Independent Fast Multipole Methods on Modern Architectures.

- a) Intended Journal: SIAM Journal on Scientific Computing
- b) Authors: Srinath Kailasa, Timo Betcke, Sarah El-Kazdadi
- c) ArXiv: 10.48550/arXiv.2408.07436
- d) Stage of Publication: Submitted
- e) Associated Thesis Chapters: 1, 3, 6
- f) Statement of Contribution
 - i. Srinath Kailasa was the lead author and responsible for the direction of this research and the preparation of the manuscript.
 - ii. Timo Betcke provided significant advisory support aid with interpretation of the results and provided critical feedback of the manuscript.
 - iii. Sarah El-Kazdadi provided expert guidance on the usage of explicit vector programming, which was critical for achieving the final results presented.

2. kiFMM-rs: A Kernel-Independent Fast Multipole Framework in Rust

- a) DOI: TODO
- b) Intended Journal: Journal of Open Source Software
- c) Authors: Srinath Kailasa
- d) Stage of Publication: Submitted
- e) Associated Thesis Chapters: 4

e-Signatures confirming that the information above is accurate

Candidate:

Date:

Supervisor/Senior Author:

Date:

Abstract

This thesis is concerned with the development of a software platform for the kernel independent Fast Multipole Method (kiFMM), a variant of the widely applied Fast Multipole Method (FMM) algorithm which finds far reaching application across computational science. Indeed, for certain dense rank-structured matrices, such as those that arise from the boundary integral formulation of elliptic Partial Differential Equations (PDEs), the FMM and its variants accelerate the computation of a matrix vector product from $O(N^2)$ to just $O(N)$ in the best case.

We demonstrate the efficacy of our software’s flexible design by contrasting implementations of a key bottleneck known as the Multipole to Local (M2L) field translation, and present a new highly optimised approach based on direct matrix compression techniques and BLAS operations, which we contrast with the current state of the art approach based on Fast Fourier Transforms (FFTs) for kiFMMs. We show that we are able to achieve highly-competitive runtimes for three dimensional problems described by the Laplace kernel with respect to the state of the art, and often faster depending on the available hardware, with a simplified approach. Our approach is well suited to the direction of development of hardware architectures, and demonstrates the importance of re-considering the design of algorithm implementations to reflect underlying hardware features, as well as the enabling power of research software for algorithm development.

The software itself is written in Rust, a modern systems programming language, with features that enable a data oriented approach to design and simple deployment to common CPU architectures. We describe our design and show how it allows us to extend our software to problems described by the Helmholtz kernel, at low frequencies, as well as in a distributed memory setting, where emphasis has been placed on retaining a simple user interface and installation suitable for non-software experts, while remaining modular enough to remain open to open-source contribution from specialists. We conclude with both single node and HPC benchmarks, demonstrating the scalability of our software as well as its state of the art performance.

Impact Statement

This thesis establishes norms and practices for developing practical implementations of the kernel independent Fast Multipole Method (kiFMM), which will be of significant utility to the developers specialising in this and related algorithms. During this research we re-visited established codes for the kiFMM, identified software construction techniques that can lead to more flexible implementations that allow users to experiment, exchange, and build upon critical algorithmic subcomponents, computational backends, and problem settings - which are often missing from competing implementations which focus achieving specific benchmarks. For example, the flexibility of the software presented in this thesis allows for the critical evaluation of key algorithmic subcomponents, such as the ‘multipole to local’ (M2L) operator which we presented in [14].

As the primary outputs are open-source software libraries [15, 13] which are embedded within existing open-source efforts, most significantly the Bempp project, with an existing user-base the software outputs of this thesis are likely to have a wide ranging impact in academia and industry influenced by the demand for these softwares. Furthermore, the adoption and promotion of Rust for this project, and within our group, establishes further the utility of this relatively new language for achieving high-performance in scientific codes, which in recent years has been the subject of growing interest in the wider high-performance scientific computing community.

Acronyms

BLAS Basic Linear Algebra Subprograms. v, 5

CLP Core Level Parallelism. 2

CPU Central Processing Unit. v, 2, 3

DLP Data Level Parallelism. 2

FFT Fast Fourier Transform. v, 5

FLOP Floating Point Operation. 5

FMM Fast Multipole Method. v, 1–5, 7–11, 23

GPU Graphics Processing Unit. 2, 3

HPC High Performance Computing. v, 4

kiFMM kernel independent Fast Multipole Method. v, 5, 6, 23

M2L Multipole to Local. v, 5, 23

MIMD Multiple Instruction Multiple Data. 2

PDE Partial Differential Equation. v, 4

SIMD Single Instruction Multiple Data. 2, 3

SIMT Single Instruction Multiple Threads. 2, 3

SISD Single Instruction Single Data. 2

TLP Thread Level Parallelism. 2

Contents

1	Introduction	1
2	Review of Fast Multipole Methods	7
2.1	Fast Multipole Methods	7
2.1.1	Intuition	7
2.1.2	Non-Adaptive Case	10
2.1.3	Adaptive Case	10
2.2	Kernel Independent Fast Multipole Methods	11
2.2.1	Motivation	11
2.2.2	Algorithm	11
2.3	Oscillatory Fast Multipole Methods	11
2.4	Related Ideas	11
2.5	Available Software	12
3	Modern Programming Environments for Science	13
3.1	Requirements for Research Software	13
3.2	Low Level or High Level? Balancing Simplicity with Performance . .	14
4	Performant Multipole To Local Field Translation for the kernel independent Fast Multipole Method	16
4.1	Forming the Multipole To Local Translation for Kernel Independent Fast Multipole Methods	16
4.2	Computational Constraints	16
4.3	Review of Acceleration Methods	16
4.4	A Direct Matrix Compression Based Acceleration Scheme	17

4.5	FFT based Acceleration Based Schemes	17
5	Software Design	18
5.1	Data Oriented Design with Rust Traits	18
5.2	FMM Software As A Framework	18
5.3	High Performance Trees	19
5.4	High Performance FMM Operators	19
5.4.1	Point to Multipole (P2M)	19
5.4.2	Multipole to Multipole (M2M) and Local to Local (L2L) . . .	20
5.4.3	Point to Point (P2P)	20
5.4.4	Multipole to Local (M2L)	20
6	The Fast Multipole Method for HPC	21
6.1	The Distributed Fast Multipole Method	21
6.2	Ghost Communication	21
6.3	Extending the Software's Design	21
7	Numerical Experiments	22
7.1	Single Node	22
7.1.1	Laplace	22
7.1.2	Helmholtz	22
7.2	Multi Node	22
8	Conclusion	23
A	Appendix	24
	Bibliography	25

Introduction

Since its introduction in the late 1980s by Greengard and Rokhlin [11], the Fast Multipole Method (FMM) has become a hallmark algorithm of scientific computing often cited as one of the ‘top 10’ algorithmic advances of the past century [7]. The problem it addresses was originally motivated by N particle simulations in which the interactions are *global* but with a strongly decaying property. Motivating examples being N particles interacting via gravity or electrostatic forces. In these cases, as well as for interactions delineated by particular interaction *kernels*, interactions between distant *clusters* of particles can be represented by truncated series expansions. This is indeed where the name for the original presentation originated, as multipole expansions derived from the fundamental solution of the Poisson equation, often called the *Laplace kernel* in FMM literature were used to form these truncated series expansions,

$$K(\mathbf{x} - \mathbf{y}) = \begin{cases} = -\frac{1}{2\pi} \log(|\mathbf{x} - \mathbf{y}|), & d = 2 \\ = \frac{1}{4\pi|\mathbf{x} - \mathbf{y}|}, & d=3 \end{cases} \quad (1.1)$$

where d is the spatial dimension. Furthermore, by using a hierarchical discretisation for the problem domain, increasingly distant interactions can be captured while still using truncated sums to express the potential due to particles contained within each subdomain in the hierarchy. With this, the FMM is able to reduce the $\mathcal{O}(N^2)$ operations required to evaluate the potentials at each of N particles due to all other particles into an algorithm requiring just $\mathcal{O}(N)$ for problems described by the Laplace kernel (1.1). The crucial advantage of the FMM is that it comes

equipped with rigorous error estimates, which guarantee exponential convergence with increasing numbers of series terms used in the truncated expansion, such that the problem could be evaluated to any desired accuracy while retaining the $\mathcal{O}(N)$ complexity bound for number of operations.

Despite being a well established algorithm, with numerous algorithmic variants and software efforts over the preceding decades, unlocking the highest available performance for practical FMM implementations remains an active area of research. Principally this can be attributed to the dramatic changes in the landscape of computing technologies in the decades since the algorithm's first introduction.

Since the end of Dennard Scaling¹ in the 2000s, hardware design and development has focussed on enhancing parallelism. Hierarchical levels of parallelism are now available to programmers representing a significant departure from the Single Instruction Single Data (SISD) systems which existed at the time of the FMM's introduction. From true hardware parallelism available via Core Level Parallelism (CLP), where multiple Central Processing Unit (CPU) cores are able to independently execute tasks or threads simultaneously in a Multiple Instruction Multiple Data (MIMD) fashion, to Thread Level Parallelism (TLP), extended by technologies such as *hyperthreading*, whereby each physical core is able to run multiple threads sharing an address space simultaneously, which are optimally scheduled by the operating system. Data Level Parallelism (DLP), whereby the same operation is applied to multiple data elements simultaneously is made available to programmers via the Single Instruction Multiple Data (SIMD) execution model represented by *vector instruction sets* available to take advantage of special hardware registers on modern Central Processing Units (CPUs), as well as via the the Single Instruction Multiple Threads (SIMT) execution model of modern Graphics Processing Units (GPUs).

The importance of simultaneously exploiting multiple levels of parallelism for achieving performance with modern software, and the increasing disparity between

¹First articulated in 1974 by Robert Dennard, Dennard scaling described how as transistors were miniaturised their power density remarkably was able to be maintained as a constant. This held true until the mid 2000s, at which point physical limits on heat dissipation and leakage current lead to power efficiency gains via miniaturisation plateauing despite the steady increase in miniaturisation described by Moore's law, marking the end of Dennard scaling. This in turn lead to the growth of multicore processors and specialised hardware accelerators, as a way to increase available computing performance without increasing power consumption.

memory bandwidth and compute resources [9], makes carefully organising memory access increasingly the key bottleneck to unlocking performance. In this context the expense of pairwise evaluations of (1.1) addressed by the FMM are increasingly less of an issue for even moderately sized problems involving hundreds of thousands of source and target particles on a single node², as the direct evaluation is embarrassingly parallel over each target particle with significant opportunity for data re-use and well suited to the SIMD or SMT execution models of modern CPUs or GPUs, respectively. Despite the $\mathcal{O}(N)$ complexity bound offered by the FMM, the operators from which it is composed contain practically significant constants and implicit non-contiguous memory accesses, making achieving practical performance challenging particularly in three dimensions.

Software efforts for Fast Multipole Methods (FMMs) were a particular focus of research activity in the 2010s, with particularly prominent examples being the ExaFMM project [4, 19], and PVFMM [17]. Indeed, software for the FMM and closely related methods have cumulatively been the recipients of three ACM Gordon Bell awards [5]. Particular recent foci have been to examine how heterogeneous architectures can be taken advantage of [17], whereby CPUs and GPUs are used in concert for data organisation and CPU bound components of the algorithm respectively, as well as taking advantage of existing parallel runtimes for achieving task-based parallelism to high effect especially in a distributed memory setting [6, 1].

The collective weakness of existing software efforts however is their brittleness. The complexity for developing high-performance software in a research setting results in softwares that are strongly adapted for a particular algorithmic approach, hardware architecture, or runtime system, often with the goal of achieving a particular benchmark or demonstrating the utility of a particular technique. Relatively little software is under active maintenance, with clear technical documentation on how performance was achieved, and fewer still open their subcomponents for extension or have downstream users. As a result, it is difficult to compare algorithmic and software optimisations taken by different implementations, and contrast their

²We demonstrate this with specific benchmarks in Section 7.1 for a selection of CPU architectures.

relative merits, as well as to see how particular approaches adapt to advances in available hardware and software.

Achieving the greatest available performance, apart from being of scientific interest for its own sake, enables larger and more detailed simulations. The FMM has been deployed as a core numerical method in the solution of elliptic Partial Differential Equations (PDEs) when formulated as boundary integral equations in combination with iterative methods, therefore acts as a crucial subcomponent in derived solvers applied to a vast array of problems in computational physics, from acoustics [12], and electromagnetics [8] to fluid dynamics [8], and biomolecular electrostatics [21, 20]. Generalised variants of the FMM have also been applied in other areas requiring fast kernel summation arising in computational statistics, such as in the widely applied Kalman filter [16], modelling Gaussian processes [2], and machine learning [10, 18]. Indeed, faster N particle kernel summations, of which the FMM is an example, have been identified as a key benchmark operation for optimisation for in HPC [3] due to their broad utility, and demonstrating the importance of performant open-source software for FMMs.

However scientific software that hopes to achieve widespread adoption must also focus on *usability* in addition to performance. Which in this context means a software that is easy to deploy on current and emerging hardware and software environments, can be interfaced by common programming languages, and is open to extension to new algorithmic approaches and implementations, while remaining highly-performant. This in turn makes the structure of the software itself an object of study.

The focus of this thesis can therefore be summarised as the development of a *platform* for developing FMMs and its variants that thrives even after the conclusion of a research project, consisting of re-usable subcomponents, and acts as a useful tool for algorithmic investigation while being capable of state of the art performance. We present this through studies into both the design as well as the application of our software in investigating the implementation of algorithmic subcomponents of the FMM.

We begin in Chapter 2 by reviewing the literature on algorithmic and software developments for FMMs, and related methods such as the \mathcal{H} and \mathcal{H}^2 matrix approaches, with a particular focus on the kernel independent Fast Multipole Method (kiFMM) a ‘black box’ variant of the FMM implemented by our software, which doesn’t rely on explicit series expansions of the fundamental solution, and only on its evaluation. We review the computational structure general of these algorithms, and identify the requirements and bottlenecks for fast implementations.

In Chapter 3 we present an investigation into the requirements of a programming platform for scientific software. Discussing the trade-off between high and low-level languages. Specifically our language of choice Rust, a modern systems-level programming language rapidly growing in popularity in science and engineering, is evaluated with respect to our previous approach that used Python, the de-facto high-level language of choice for high-level scientific computing and rapidly developing features that enable high-performance, adapting material first presented in [15].

In Chapter 4 we present a rigorous application of our software, in which we investigate optimisations for a crucial Multipole to Local (M2L) operation for kiFMMs, a critical bottleneck in FMMs due to its requirement for significant data re-organisation due to the implicit non-contiguous data accesses required by this operation. This operation is of convolution type, and current state of the art approaches for the the M2L are based on Fast Fourier Transforms (FFTs) with an $\mathcal{O}(N \log N)$ complexity bound, requiring careful explicit vector programming to achieve practical performance. We find that despite this we are able to construct a highly competitive scheme based on direct matrix compression techniques and Basic Linear Algebra Subprograms (BLAS) operations due to the superior memory re-use of the approach, despite it requiring a greater number of Floating Point Operations (FLOPs). Indeed, our scheme’s reliance on BLAS operations, and comparatively simple algorithmic structure and implementation, make it well suited to direction of development of computer hardware, and depending on the underlying hardware can even offer faster performance than the current state of the art approach. The material in this chapter is largely adapted from that first presented in [14].

Chapter 5 describes in detail the engineering approach of our software, particularly the employment of Rust’s ‘trait’ system, as well as specific implementation details of the kiFMMs operators, and data structures. Chapter 6 describes the extension of our software to distributed memory systems, and details communication reducing schemes for exchanging ghost information. Chapter 7 contains numerical experiments with our software in a single node as well as HPC setting, including a study of the applicability of our software to oscillatory problems described by the time independent Helmholtz equation at low to moderate wavenumbers. We conclude with a reflection on our results and suggestions for future avenues of investigation in Chapter 8.

Review of Fast Multipole Methods

Portions of the discussion in Sections 2.1 and 2.2 of this chapter are adapted from the material first presented in [14]

2.1 Fast Multipole Methods

2.1.1 Intuition

As in the original presentation, we use the case of evaluating electrostatic potentials to motivate the FMM. Consider the electric field, \mathbf{E} due to a static charge distribution $q(\mathbf{y})$ which is supported over some finite domain $\mathbf{y} \in \Omega \subset \mathbb{R}^d$. It can be defined in terms of a scalar potential ϕ .

$$\mathbf{E} = -\nabla\phi$$

which itself can be seen to satisfy Poisson's equation,

$$\begin{cases} -\Delta\phi(\mathbf{x}) = q(\mathbf{x}), & \text{for } \mathbf{x} \in \mathbb{R}^d \\ \lim_{|\mathbf{x}| \rightarrow \infty} \phi(\mathbf{x}) = 0 \end{cases}$$

where $d = 2, 3$ in problems of interest.

We can write the evaluation of the potential at a point \mathbf{x} as a convolution of the source with the fundamental solution of the Poisson equation,

such that,




Figure 2.1: A set of source and target particle cluster, where the width of each cluster is significantly less than the distance separating them, $d_s, d_t \ll D$.

$$\phi(\mathbf{x}) = \int_{\mathbb{R}^d} K(\mathbf{x} - \mathbf{y})q(\mathbf{y})d\mathbf{y}, \quad \mathbf{x} \in \mathbb{R}^d \quad (2.1)$$

Under an appropriate discretisation, where care is taken to appropriately handle the singularity in the Laplace kernel (1.1), we see that this integral corresponds to a matrix vector multiplication, where the matrix is *dense*, i.e. it consists only of non-zero entries.

As we are principally concerned with the simpler problem of evaluating the potential due to a discrete charge distribution, with N charges we can replace $q(\mathbf{y})$ with $\{q(\mathbf{y}_j)\}_{j=1}^N$ associated with *source particles* $\{\mathbf{y}_j\}_{j=1}^N \in \mathbb{R}^d$, the integral for potential evaluated at M *target particles*, $\{\mathbf{x}_i\}_{i=1}^M \in \mathbb{R}^d$ becomes a discrete sum,

$$\phi(\mathbf{x}_i) = \sum_{j=1}^N K(\mathbf{x}_i - \mathbf{y}_j)q(\mathbf{y}_j), \quad i = 1, \dots, M \quad (2.2)$$

where we can handle the singularity by setting,

$$K(\mathbf{x}_i - \mathbf{y}_j) = \begin{cases} 0, & \mathbf{x}_i = \mathbf{y}_j \\ K(\mathbf{x}_i - \mathbf{y}_j), & \text{otherwise} \end{cases} \quad (2.3)$$

We see that the sum (2.2) corresponds to a dense matrix vector multiplication,

$$\phi = Kq \quad (2.4)$$

Naively computed this requires $\mathcal{O}(MN)$ operations, where in general the source and target particles may correspond to the same set. The FMM relies on a *degenerate* approximation of the interaction kernel when subsets, or *clusters*, of source and

target particles are sufficiently separated as sketched in Figure 2.1. Following the discussion in [14] the sum (2.2) can be written as,

$$\phi(\mathbf{x}_i) \approx \sum_{p=1}^P \sum_{j=1}^N A_p(\mathbf{x}_i) B_p(\mathbf{y}_j) q(\mathbf{y}_j), \quad i = 1, \dots, M \quad (2.5)$$

where we call P the expansion order, taken such that $P \ll N$, $P \ll M$. The functions A_p and B_p are defined by the approximation scheme used by a particular approach for the FMM, in the original presentation the calculation,

$$\hat{q}_p = \sum_{j=1}^N B_p(\mathbf{y}_j) q(\mathbf{y}_j) \quad (2.6)$$

Corresponded to the coefficients of an order P multipole expansion due to the source charges. Following which the potential is approximated by,

$$\phi(\mathbf{x}_i) \approx \sum_{p=1}^P A_p(\mathbf{x}_i) \hat{q}_p, \quad i = 1, \dots, M \quad (2.7)$$

at the target particles. The approximation of the potential with this scheme can be seen to require $\mathcal{O}(P(M + N))$ operations. The accuracy of this approximation scheme, and the error bounds provided by the FMM, depends on the distance between the source and target clusters remaining large relative to their width. This condition is often referred to as an *admissibility condition* in the FMM literature. FMMs therefore split the sum (2.2) into *near* and *far* components when considering arbitrary clusters of source and target particles,

$$\phi(\mathbf{x}_i) = \sum_{\mathbf{y}_j \in \text{Near}(\mathbf{x}_i)} K(\mathbf{x}_i, \mathbf{y}_j) q(\mathbf{y}_j) + \sum_{\mathbf{y}_j \in \text{Far}(\mathbf{x}_i)} K(\mathbf{x}_i, \mathbf{y}_j) q(\mathbf{y}_j), \quad i = 1, \dots, M \quad (2.8)$$

In cases where a source and target cluster can be considered *admissible*, i.e. the source cluster is considered in the *far field* of the target cluster such that each $\mathbf{y}_j \in \text{Far}(\mathbf{x}_j)$, we apply the approximation (2.5). However, when a source and target cluster are *inadmissible*, such that the source cluster is considered in the *near field* of a target cluster such that each $\mathbf{y}_j \in \text{Near}(\mathbf{x}_j)$ we are left to evaluate the sum

directly via (2.2).

In order to achieve its $\mathcal{O}(N)$ complexity the FMM is structured to reduce to a minimum the number of sums evaluated between inadmissible clusters.

2.1.2 Non-Adaptive Case

- How are analytical expansions defined, 2D Laplace example, note and references on 3D laplace example. - private note on how these derivations are found.

- Note on admissibility, how it defines a broad class of FMM matrices, table of related matrix types with different rank structure and admissability criterion

This is achieved with a hierarchical discretisation of the problem domain, often a *quadtree* in two dimensions and correspondingly an *octree* in three dimensions.

- trees define admissability from geometry for FMM, note on alternative approaches such as ORB.

These data structures are generated by creating a bounding box that covers the source and target particles, which without loss of generality may correspond to the same set. This box is then recursively sub-divided into *child boxes* of equal size.

- Define interaction lists.
- Define all operators.
- Give an algorithm sketch.
- Give complexity sketch.

2.1.3 Adaptive Case

- Types of adaptivity available (weak, strong) requirements on interaction lists.
 - Define new interaction lists, and new operators
 - Give complexity sketch.

2.2 Kernel Independent Fast Multipole Methods

2.2.1 Motivation

The decades since its original presentation have seen the FMM extended with related ideas which principally differ in how they represent interactions between distant clusters of particles typically referred to as *algebraic* FMMs.

- Review of the KiFMM and variants. Black Box FMM, Analytical FMM, Data Driven Techniques.

- Motivation for use from a software engineering and computational performance perspective.

2.2.2 Algorithm

- Approximation scheme via method of fundamental solutions, why do we expect certain convergence?

2.3 Oscillatory Fast Multipole Methods

- helmholtz fmm

- high frequency helmholtz FMM, out of scope of the thesis but can mention that they exist. Especially in the context of kiFMM approaches, what is the key difference? When does it apply (the directional low rank condition), where is there an additional loop? Why might this result in greater complexity

2.4 Related Ideas

- \mathcal{H} Matrix and \mathcal{H}^2 matrices, and wider setting of the FMM and related problems.
- Abduljabbar thesis contains a nice summary I can read.

2.5 Available Software

- What software exists, and which approaches do they use?
 - What are they optimised for, and what kind of performance do they promise?
 - What are the trade-offs of each software
 - Hardware targetted by each available software, what's missing?
 - What's available, and what are the shortcomings?
 - And then

Modern Programming Environments for Science

The discussion in this chapter, including figures and diagrams, is adapted from the material first presented in [15].

3.1 Requirements for Research Software

- Requirements and constraints on research software development.

- As an example FMM softwares used in recent benchmark studies (ExaFMM variants, PVFMM) have been constructed during the course of doctoral or post-doctoral projects. This entails a significant ‘key man’ risk, in which when the project owner completes their course of research the project enters a decay state and is no longer actively maintained and developed. New developers, unfamiliar with the code bases which can grow to thousands of lines of code, and often written without reference to standard software engineering paradigms for designing and managing large code bases (continuous integration, software diagrams, and simple decoupled interfaces) will find it challenging to build upon existing advances, and resort to developing new code-bases from scratch, rediscovering implementation details that are often critical in achieving practical performance.

- In seeking to avoid this cycle we envisioned a project built in Python, which maximises the maintainability of a project due to its simple syntax and language construction. New developers who, as a standard, are often educated in Python in the natural sciences and engineering, will hopefully be familiar with the language in

order to gain productivity as fast as possible. However, as we demonstrated in our paper ... This itself imposes significant constraints on performance, which is balanced by the 'usability' of the language, making it just as challenging as developing a complex code in a compile language.

- Modern compiled languages offer tools that enable developer productivity. Examples include Go, Rust, ...

- The complexity of methods leads to complex code surface areas which are difficult to maintain especially in an academic setting with few resources for professional software engineering practice.

- The diversity of hardware and software backends leads to increasing difficulty for projects to experiment with and incorporate computational advances.

- Hardware and software complexity, and gap between a one-off coding project and extensible maintainable software tooling.

- Review developments in computer hardware and software that make this easier to be more productive, but also more challenging to wrap together over time.

- Emerging and future trends, exemplified by the step change in compiled languages in the new generation and the interest in Rust and similar languages. The mojo project and what this says about the future.

3.2 Low Level or High Level? Balancing Simplicity with Performance

- Summary of Python paper results, in summary complex algorithms necessitate complex code in order to achieve performance - specifically the requirement for programmers to be in charge of memory and for hot sections manually vectorise etc. Writing everything in a high-level language obfuscates the application code from the sections critical to performance

- Review of why this was thought to be a good idea, and why it might be worth trying again in the future.

- What problems does this paper address, wrt to the literature?

- Brief review of motivation and reasoning behind Rust, and which features we take advantage of
- Build system, paragraph explaining it and contrast with build systems wrt to competitors.
- Trait system, though shares similarities with C++21 concepts, a part of the type systems - organisation of shared behaviour without loss of performance, bottom up organisation.
- Why the two language problem isn't really a problem with modern compiled languages, note on the successful projects that manage this and how effective just writing C interfaces can be, if using another language.

Performant Multipole To Local Field Translation for the kernel independent Fast Multipole Method

The discussion in this chapter, including figures and diagrams, is adapted from the material first presented in [14]

4.1 Forming the Multipole To Local Translation for Kernel Independent Fast Multipole Methods

4.2 Computational Constraints

- Exactly why is this hard to evaluate? - Maybe complexity estimates of memory accesses. - Layout discussion.

4.3 Review of Acceleration Methods

- Full literature review of past approaches - brief summary of analytical approaches (1-2 paragraphs) - Where past efforts have been focussed, and why? (Original paper dismissed direct matrix compression) - Why this may be a good idea now (randomised methods, available hardware)

- use this section to introduce idea of transfer vectors, reflection and rotational symetry - how this is achieved in practice (i.e. what computations are needed, not the implementation details)

Fast Multipole Method as a Matrix-Free Hierarchical Low-Rank Approximation

Rio Yokota, Huda Ibeid, David Keyes

4.4 A Direct Matrix Compression Based Acceleration Scheme

- Why might this be preferred, or advantageous, what are its constraints - Why it is counterintuitive

4.5 FFT based Acceleration Based Schemes

- Explanation of the method, and why it was able to achieve high performance.
- Why this may not be completely appropriate, low arithmetic intensity (maybe estimate?)

- Both algorithmic and computational

- Review of approaches, success and failures, and what works in the context of modern software and hardware systems.

- Can include section reviewing PVFMM approach

- Approaches for BLAS based field translation in some more detail than in the paper.

- Our approach, and why it works with reference to the software and hardware currently available.

- Articulate the significance of this result

- Why our approach is sustainable given long term trends in hardware and software.

- Why might it be useful for Helmholtz FMM ...

- What are important trends, and what have we actually done.

Software Design

5.1 Data Oriented Design with Rust Traits

- Motivation, and review, DOD book.
 - How do traits enable data oriented design.
 - Overview of the design of the software.
 - Diagram for principal traits and how they link together in the final software.
 - Why is this good for the future? Well, it leaves open extension to other approaches for any individual subcomponent.
- An example of this is the genericity over data type, kernel implementation, and field translation method, with a space for the kind of tree data structure.

5.2 FMM Software As A Framework

- Want to encourage as much code re-use as possible.
 - The re-implementation of critical subcomponents should be avoided. A step towards this is the development low-level C interfaces which enable the construction of higher level interfaces in compatible languages.
 - We've made a start to this with a low-level interface to the principal API of the FMM software.
 - We also want to be able to deploy on as wide a range of target hardware as possible, and leave open extension to future systems, enabled by design, referencing diagram.
 - High level diagram of how software components fit together

- Explain the diagrams (will need zoomed in diagrams for things like M2L data)
- Decouple implementation from
- Code generation for multiple targets enabled by Rust's llvm based compiler.
- C ABI as a compatibility layer to other projects, success with this in developing Python wrappers and integration with NGBEM
- Flexible backends enabled by RLST package for BLAS and Lapack.

5.3 High Performance Trees

- Exactly how are Morton encodings done, and what are the drawbacks and alternatives.
 - ORB (how does it work)
 - Tree Construction approach and algorithms - Morton encoding via lookup tables - neighbour finding - interaction list construction (fast)
 - What did we end up doing, and what is the justification for these being good enough.

Important implementation details - construction of interaction lists, neighbour finding. - construction of Morton encodings. - trade-offs of approach in shared and distributed memory - e.g. adaptive vs weakly adaptive trees. - problems with load balancing approach etc - justification - simplicity/works vs complex/private

5.4 High Performance FMM Operators

- FMM data flow, mapping its tree structure to a data flow diagram.
 - Go through each operator of kiFMM and how it's been optimised, at a high level for M2L operators.

5.4.1 Point to Multipole (P2M)

- Formulation in a blocked manner

5.4.2 Multipole to Multipole (M2M) and Local to Local (L2L)

- Formulation as BLAS3

5.4.3 Point to Point (P2P)

- Computational challenge (SIMD/SIMD), and why this is the easiest to optimise.

5.4.4 Multipole to Local (M2L)

- Description of computational challenge due to memory layout proscribed by Morton encoding.

- Formulation of the problem, and required precomputations.
- Data structure required to be flexible over this.

How future FMMs may look, shallow trees with large P2P.

The Fast Multipole Method for HPC

6.1 The Distributed Fast Multipole Method

6.2 Ghost Communication

6.3 Extending the Software's Design

- Focus is on the maximal reduction in communication.
 - what communication can and cannot be avoided?
 - How the local/global split in terms of tree gives rise to optimal communication scheme.
 - What simplifying assumptions can we take for most pre-exascale systems?
 - Avoid sorting of Morton keys/point data, the local/global split gives us a way to statically partition tree across available resources - simplifying assumption if work with $n_{cpu} = \text{pow}(8)$. - Not restricted to this, but makes threading simpler for local FMMs.
 - Optimal implementation of MPI primitives for common data sizes.
 - What will probably not work approaching exascale? - the gather operation over all processes required for multiple steps of this algorithm - ghost exchange, multipoles at local root on nominated processor. How can these problems be addressed? Do they even matter for the problem sizes we're concerned with?

Numerical Experiments

7.1 Single Node

7.1.1 Laplace

7.1.2 Helmholtz

7.2 Multi Node

Conclusion

In this thesis we have presented progress on the development and design of a software framework for kernel independent Fast Multipole methods. We've documented outputs towards the broader goal of a sustainable framework which can be extended, with re-usable subcomponents. This research performed necessitated a significant investigation into the optimal programming environment for high-performance scientific computing that enabled high productivity within the constraints of academic software development. The resultant software enabled a new investigation of the critical M2L field translation operation, a key bottleneck in the kiFMM algorithm, and the development of a highly competitive approach well suited to emerging trends in computer hardware.

Due to the high-performance the FMM operator kernels for both the Laplace and low-frequency Helmholtz kernels demonstrated in this work as well as the creation of trees, we are in a good position to extend our software to a distributed setting. The decoupling of operator kernels from their implementation via the design of our software also enables future extensions to a heterogenous platforms in which batch BLAS for the M2L, and SIMT for the P2P operations

- Direction of travel of hardware, how this could effect exactly how the algorithm is set up
- specifically handling as much as possible on a GPU, data organisation in a unified memory context will be cheap.

Appendix

Bibliography

- [1] Emmanuel Agullo et al. “Task-based FMM for multicore architectures”. In: *SIAM Journal on Scientific Computing* 36.1 (2014), pp. C66–C93.
- [2] Sivaram Ambikasaran et al. “Fast direct methods for Gaussian processes and the analysis of NASA Kepler mission data”. In: *arXiv preprint arXiv:1403.6015* (2014).
- [3] Krste Asanovic et al. “The landscape of parallel computing research: A view from berkeley”. In: (2006).
- [4] Lorena A Barba and Rio Yokota. “ExaFMM: An open source library for Fast Multipole Methods aimed towards Exascale systems”. In: *Boston: Boston University. Retrieved from barbagroup: <http://barbagroup.bu.edu>* (2011).
- [5] Gordon Bell et al. “A look back on 30 years of the Gordon Bell Prize”. In: *The International Journal of High Performance Computing Applications* 31.6 (2017), pp. 469–484.
- [6] B  renger Bramas. “TBFMM: A C++ generic and parallel fast multipole method library”. In: *Journal of Open Source Software* 5.56 (2020), p. 2444.
- [7] Barry A Cipra. “The best of the 20th century: Editors name top 10 algorithms”. In: *SIAM news* 33.4 (2000), pp. 1–2.
- [8] Eric Darve and Pascal Hav  . “A fast multipole method for Maxwell equations stable at all frequencies”. In: *Philosophical Transactions of the Royal Society of London. Series A: Mathematical, Physical and Engineering Sciences* 362.1816 (2004), pp. 603–628.

- [9] Jack Dongarra et al. “With extreme computing, the rules have changed”. In: *Computing in Science & Engineering* 19.3 (2017), pp. 52–62.
- [10] Alexander Gray and Andrew Moore. “N-Body problems in statistical learning”. In: *Advances in neural information processing systems* 13 (2000).
- [11] Leslie Greengard and Vladimir Rokhlin. “A fast algorithm for particle simulations”. In: *Journal of computational physics* 73.2 (1987), pp. 325–348.
- [12] Sijia Hao, Per-Gunnar Martinsson, and Patrick Young. “An efficient and highly accurate solver for multi-body acoustic scattering problems involving rotationally symmetric scatterers”. In: *Computers & Mathematics with Applications* 69.4 (2015), pp. 304–318.
- [13] Srinath Kailasa. “kifmm-rs: A Kernel-Independent Fast Multipole Framework in Rust”. Submitted to Journal of Open Source Software. 2024.
- [14] Srinath Kailasa, Timo Betcke, and Sarah El Kazdadi. *M2L Translation Operators for Kernel Independent Fast Multipole Methods on Modern Architectures*. 2024. arXiv: 2408.07436 [cs.CE]. URL: <https://arxiv.org/abs/2408.07436>.
- [15] Srinath Kailasa et al. “PyExaFMM: an exercise in designing high-performance software with Python and Numba”. In: *Computing in Science & Engineering* 24.5 (2022), pp. 77–84.
- [16] Judith Yue Li et al. “A Kalman filter powered by-matrices for quasi-continuous data assimilation problems”. In: *Water Resources Research* 50.5 (2014), pp. 3734–3749.
- [17] Dhairya Malhotra and George Biros. “PVFMM: A parallel kernel independent FMM for particle and volume potentials”. In: *Communications in Computational Physics* 18.3 (2015), pp. 808–830.
- [18] William B March and George Biros. “Far-field compression for fast kernel summation methods in high dimensions”. In: *Applied and Computational Harmonic Analysis* 43.1 (2017), pp. 39–75.

- [19] Tingyu Wang, Rio Yokota, and Lorena A Barba. “ExaFMM: a high-performance fast multipole method library with C++ and Python interfaces”. In: *Journal of Open Source Software* 6.61 (2021), p. 3145.
- [20] Tingyu Wang et al. “High-productivity, high-performance workflow for virus-scale electrostatic simulations with Bempp-Exafmm”. In: *arXiv preprint arXiv:2103.01048* (2021).
- [21] Rio Yokota et al. “Biomolecular electrostatics using a fast multipole BEM on up to 512 GPUs and a billion unknowns”. In: *Computer Physics Communications* 182.6 (2011), pp. 1272–1283.